



HashiCorp  
**Terraform**

# Scaling Infrastructure Automation

# Course Agenda

## Terraform 201 – Scaling Infrastructure Automation

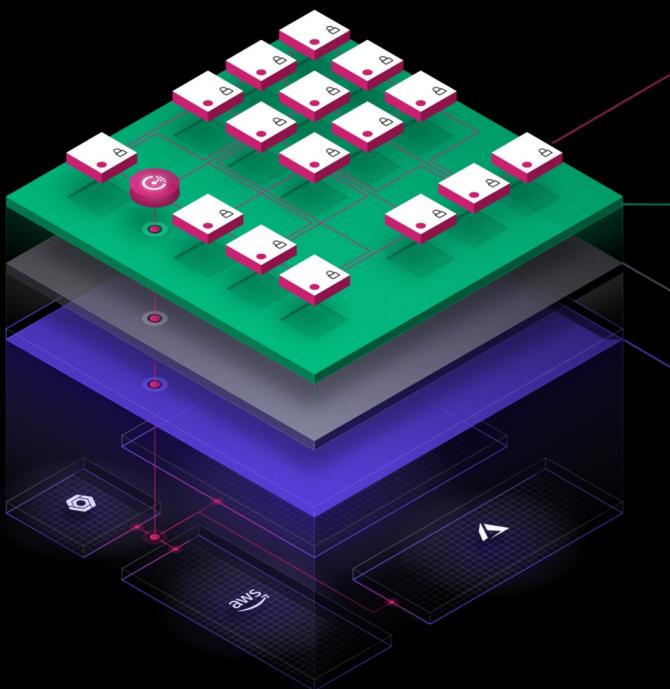
Course Agenda:

- . Terraform Cloud/Enterprise Overview
- . Version Control
- . Remote State Management
- . Variable Management
- . Data Blocks
- . Conditional Logic
- . Validation and Suppression
- . Local Values
- . For-Each Blocks
- . Dynamic Blocks
- . VCS Integration and Workflow
- . Private Module Registry
- . Null Resource and Remote State Data



# HashiCorp Ecosystem

Our open-source offerings for managing infrastructure.



- **Connect**

Consul - runtime service discovery and distributed key-value stores

- **Run**

Nomad - distributed application scheduler

- **Secure**

Vault - manages secrets across dev, ops, prod, etc.

- **Provision**

Terraform - provisions infrastructure resources



## Terraform's Goals

- Infrastructure as Code
- Modern data center support
- Safe, predictable operation
- Tech Agnostic
- Manage anything with an API



- **Connect**

Consul - runtime service discovery and distributed key-value stores

- **Run**

Nomad - distributed application scheduler

- **Secure**

Vault - manages secrets across dev, ops, prod, etc.

- **Provision**

Terraform - provisions infrastructure resources

# Introductions



- Background
- What would you like to get out of the class?
- AWS Experience? Azure? GCP?
- Coolest thing you have done with Terraform to date?
- What Time Zone are you joining from?

# Overview



## We will cover

- Tools for teams and scale
- Terraform Cloud
- Examples use AWS
- Q&A and your requests

## We will NOT cover

- Terraform Enterprise (OnPrem)
- Terraform Enterprise API
- Source code management tools
- Configuration management tools
- Every cloud

# What You'll Learn



- ✓ Use collaborative features of Terraform
- ✓ Organize Terraform projects for large scale infrastructure installations
- ✓ Use state
- ✓ Become confident with supporting tools



---

Terraform 201

# Review

# Key Terraform Syntax & Features



- ➊ Summary of providers, resources, variables, outputs
- ➋ Summary of modules
  
- ➌ Tips for upgrading commands and providers

# Terraform's Goals

---

Unify the view of resources using infrastructure as code

Support the modern data center (IaaS, PaaS, SaaS)

Expose a way for individuals and teams to safely and predictably change infrastructure

Provide a workflow that is technology agnostic

Manage anything with an API

# Terraform vs. Other Tools

---

Provides a high-level abstraction of infrastructure  
(IaC)

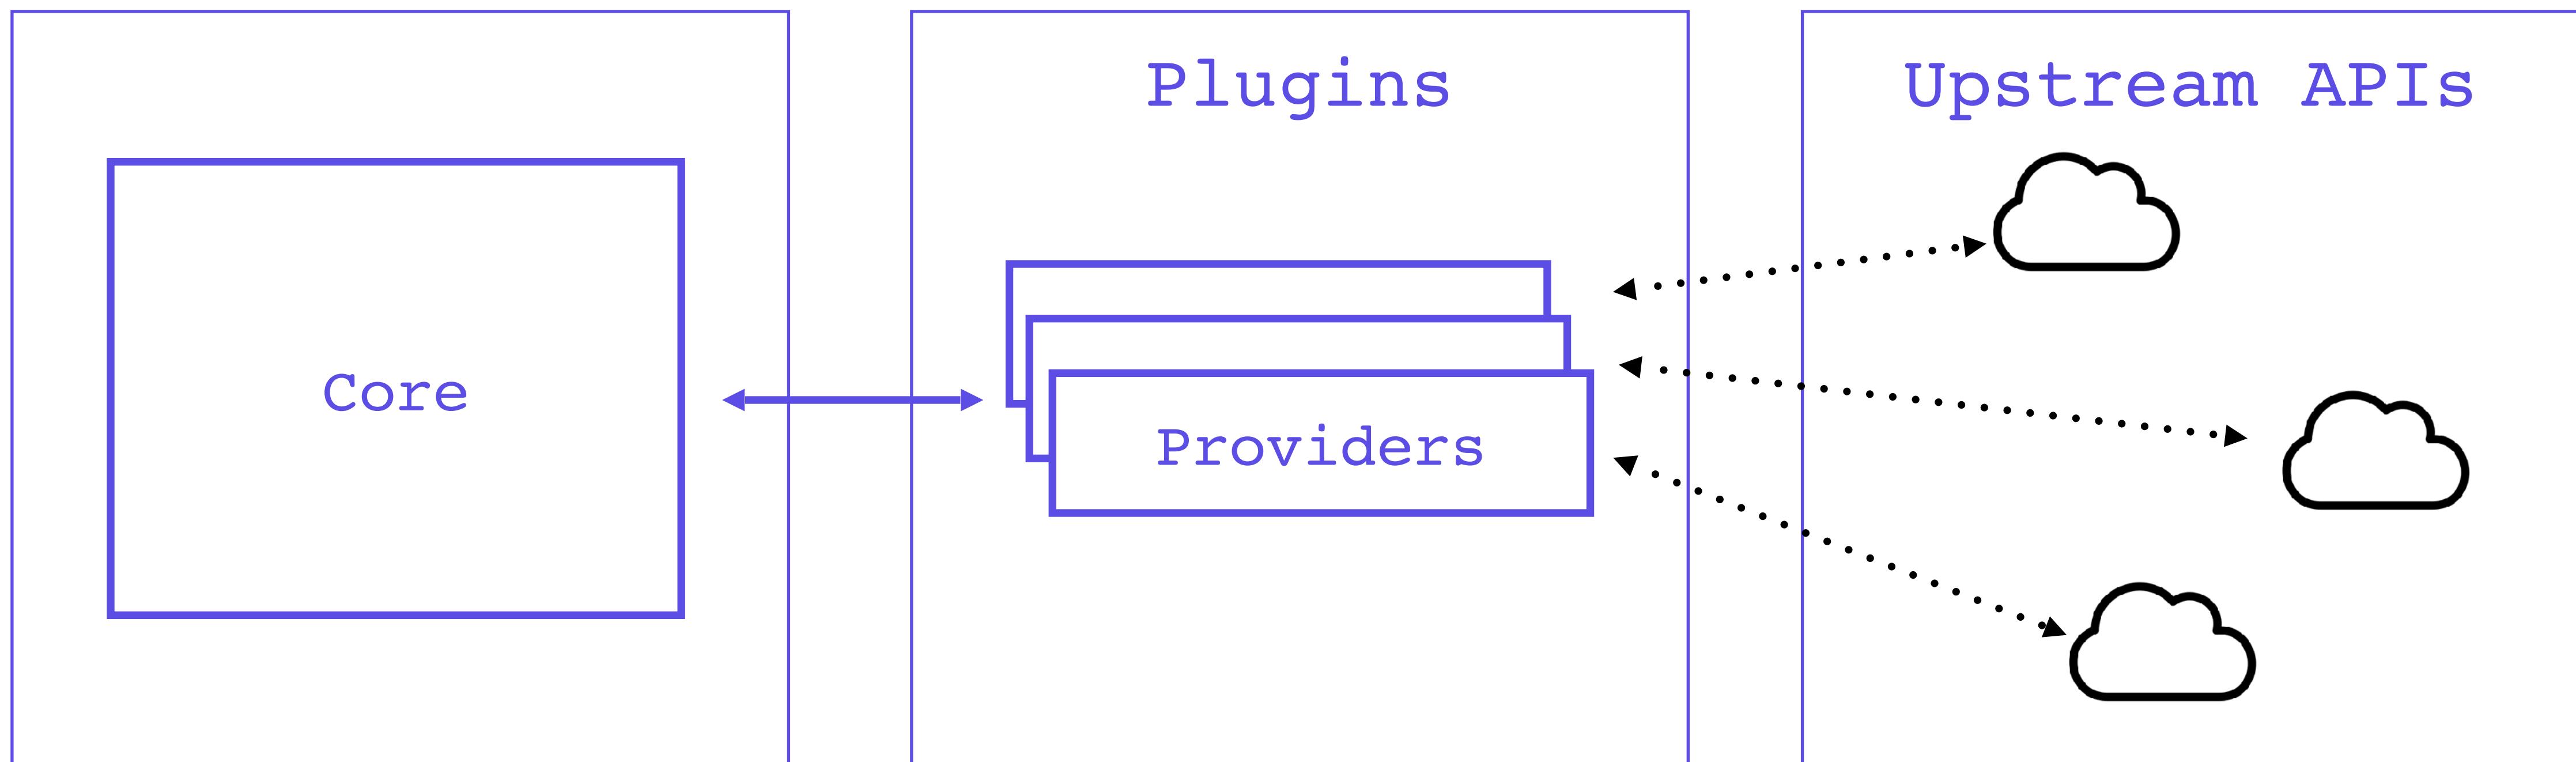
Allows for composition and combination

Supports parallel management of resources  
(graph, fast)

Separates planning from execution (dry-run)

# Terraform's Internals: Structure

---



# Configuration Syntax

Configurations are written in HashiCorp Configuration Language (HCL)

HCL is designed to strike a balance between human-readable and machine-parsable

Terraform can also read configuration in JSON, but we recommend using HCL, except when the configuration is machine-generated

# Top Level Keywords

provider

variable

resource

output

module

data

terraform

locals

# Terraform's Internals: Core Concepts

---

**Config:** Target Reality

**State:** Current Reality

**Diff:** {Config - State}

**Plan:** Presents Diff

**Apply:** Resolves Diff

# Terraform Syntax



Providers - Cloud, VCS, Monitoring, etc.

Resources - EC2 Instances, IAM Users, etc.

Variables - Input for tagging or defining resources.

Outputs - Used for record keeping or as a variable  
to pass into other modules.

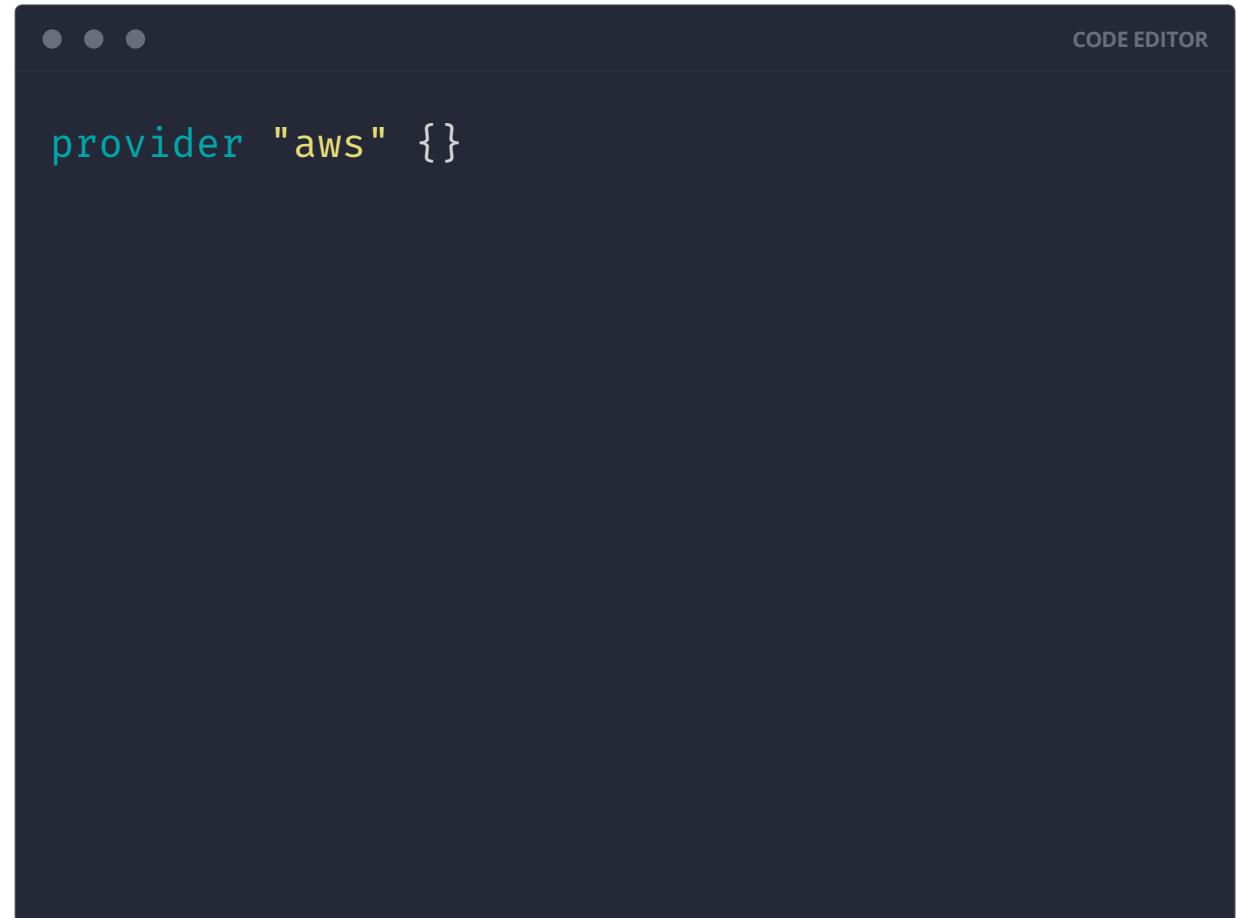
Modules - Segmented configurations for easy  
repeatability.



---

# Providers

Declaring plugins for  
Terraform to use in  
building infrastructure



A screenshot of a dark-themed code editor window. At the top right, the text "CODE EDITOR" is visible. In the center of the editor area, there is a single line of code: "provider "aws" {}". Above the code editor, there are three small circular dots, likely indicating a navigation or search feature.

# Secure Credentials with ENV Variables or Cloud Role



For flexibility and to nearly eliminate the possibility of accidentally committing cloud credentials to source code control systems, consider storing cloud credentials in environment variables. Or use cloud role.

If they are named correctly, the appropriate Terraform provider will discover and use them.

```
export AWS_ACCESS_KEY_ID="AAAAAAA"  
export AWS_SECRET_ACCESS_KEY="XXXXXXXXXX"  
export AWS_DEFAULT_REGION="us-west-2"
```



# Securely Sourcing Credentials

With an AWS CLI profile, we can load credential profiles into our configuration



A dark-themed terminal window titled "TERMINAL". At the top, there are three small circular icons. Below the title, the command "\$ source ~/.config/envs/aws" is displayed in white text. The rest of the terminal window is blank black space.



## Important :

### **Providers MUST be explicitly upgraded**

We value stability and we respect your control over your environment.

By default, Terraform will lock on to the first version available when you first run init on a configuration.

To upgrade, you must explicitly state a version in your configuration, or run this command:

```
terraform init -upgrade
```



# Provider Versioning

You can explicitly declare the version number of the provider you want to use.

```
provider "aws" {  
    version = ">= 1.19.0"  
}
```

CODE EDITOR



## Resources

The resource stanza declares what infrastructure

Terraform will build and the attributes of those resources

```
provider "aws" {}

resource "aws_instance" "web" {
    ami           = "ami-e5d9439a"
    instance_type = "t2.nano"
}
```

CODE EDITOR



# Variables

The variable stanza defines the parameters of data to be interpreted into our configuration later

```
variable "ami" {
  description = "The Amazon Machine Image ID"
  default     = "ami-12345"
}

provider "aws" {}

resource "aws_instance" "web" {
  ami           = var.ami
  instance_type = "t2.nano"
}
```



## More variables!

Variables should be moved to separate "variables.tf" files.

Terraform can parse map and list types.

```
variable "images" {  
    type = "map"  
    default = {  
        us-east-1 = "image-1234"  
        us-west-2 = "image-4567"  
    }  
}  
  
variable "zones" {  
    type = "list"  
    default = ["us-east-1a", "us-east-1b"]  
}
```



# Outputs

Outputs can render data from the state file for later use.

Should also be moved to a separate “output.tf” file for maintainability

```
provider "aws" {}

resource "aws_instance" "web" {
    ami           = var.images
    instance_type = "t2.nano"
}

output "public_ip" {
    value = aws_instance.web.public_ip
}
```



# Directory Layout

A common Terraform  
folder structure

TERMINAL

```
my-infrastructure
├── main.tf
├── output.tf
├── terraform.tfvars
├── terraform.tfstate
└── variables.tf
```



## Using modules

Modules are referenced with the local path.

```
module "vpc" {  
  source = "../vpc_module"  
}
```

CODE EDITOR



## Modules:

### File Structure

This is an example of a file structure for a module nested within our current directory.

```
TERMINAL
$ tree my-infrastructure/
.
├── main.tf
├── output.tf
├── terraform.tfvars
├── terraform.tfstate
└── variables.tf
    └── vpc_module/
        ├── main.tf
        ├── output.tf
        └── variables.tf
```



# Built-in Functions

There are several functions for working with lists, files, and more

```
# Useful built-in functions  
basename(path)  
cidrhost(prefix, hostnum)  
concat(list1, list2, ...)  
file(path)  
join(separator, list)  
  
my-list[7]
```

CODE EDITOR

<https://www.terraform.io/docs/configuration/functions.html>

# Versioning

# Versioning

Terraform consults version constraints to determine whether it has acceptable versions of itself, any required provider plugins, and any required modules.

For plugins and modules, it will use the newest installed version that meets the applicable constraints.

If Terraform doesn't have an acceptable version of a required plugin or module, it will attempt to download the newest version that meets the applicable constraints.

# Versioning

If Terraform isn't able to obtain acceptable versions of external dependencies, or if it doesn't have an acceptable version of itself, it won't proceed with any plans, applies, or state manipulation actions.

Both the root module and any child module can constrain the acceptable versions of Terraform and any providers they use. Terraform considers these constraints equal, and will only proceed if all of them can be met.

# Version Syntax

A version constraint is a string literal containing one or more conditions, which are separated by commas.

Each condition consists of an operator and a version number.

# Version Syntax

## Valid Operators:

- = : Allows only one exact version number. Cannot be combined with other conditions.
- != : Excludes an exact version number.
- >, >=, <, <=: Comparisons against a specified version, allowing versions for which the comparison is true.
- ~>: Allows only the rightmost version component to increment. For example, to allow new patch releases within a specific minor release, use the full version number: ~> 1.0.4 will allow installation of 1.0.5 and 1.0.10 but not 1.1.0.

```
● ● ●
```

```
terraform {
  required_version = ">= 1.1.0"
  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~3.0"
    }
  }
}
```

# Dependency Lock

Terraform configurations must declare which providers they require so that Terraform can install and use them.

Terraform will record the provider versions in a dependency lock file to ensure that others using this configuration will utilize the same Terraform and provider versions.

This file by default is saved as a *.terraform.lock.hcl* file

```
provider "registry.terraform.io/hashicorp/aws" {
    version      = "3.73.0"
    constraints = "~> 3.0"
    hashes = [
        "h1:fpZ14qQnn+uE002Z0lBFHgty480l8I0wd+ewxZ4z3zc=",
        "zh:0483ca802ddb0ae4f73144b4357ba72242c6e2641aeb460b1aa9a6f6965464b0",
        "zh:274712214ebeb0c1269cbc468e5705bb5741dc45b05c05e9793ca97f22a1baa1",
        "zh:3c6bd97a2ca809469ae38f6893348386c476cb3065b120b785353c1507401adf",
        "zh:53dd41a9aed9860adbbeeb71a23e4f8195c656fd15a02c90fa2d302a5f577d8c",
        "zh:65c639c547b97bc880fd83e65511c0f4bbfc91b63cada3b8c0d5776444221700",
        "zh:a2769e19137ff480c1dd3e4f248e832df90fb6930a22c66264d9793895161714",
        "zh:a5897a99332cc0071e46a71359b86a8e53ab09c1453e94cd7cf45a0b577ff590",
        "zh:bdc2353642d16d8e2437a9015cd4216a1772be9736645cc17d1a197480e2b5b7",
        "zh:cbeace1deae938f6c0aca3734e6088f3633ca09611aff701c15cb6d42f2b918a",
        "zh:d33ca19012aab98cc03fdecc0bd5ce56e28f61a1dfbb2eea88e89487de7fb3",
        "zh:d548b29a864b0687e85e8a993f208e25e3ecc40fcc5b671e1985754b32fdd658",
    ]
}
```

# Provider Upgrades

---

By default, if a `.terraform.lock.hcl` file exists within a Terraform working directory, the provider versions specified in this file will be used.

This lock file helps to ensure that runs accross teams will be consistent.

To upgrade provider versions we update our configuration definition to the desired provider version and running an upgrade.

```
> terraform init -upgrade
```

```
terraform {
    required_version = "1.1.0"

    required_providers = {

        aws = {
            version = "~>3.0"
            source = "hashicorp/aws"
        }
    }
}
```

---

## EXERCISE

## Lab 1: Getting Started

Duration: 15 minutes

Goal: Familiarize yourself with your workstation  
and deploy a basic configuration

---

## EXERCISE

## Lab 2: Versions

Duration: 10 minutes

Goal: Control the version of Terraform and AWS provider the code is compatible with

# Terraform OSS and TFC/TFE

# Terraform Ecosystem

---

Terraform CLI	Terraform Cloud	Terraform Enterprise
Command line tool	SaaS app	On premise/in your cloud
Infrastructure as Code	Free for individuals and teams	Workspace & team management
Open Source	Remote state	Governance and policy features
150+ Providers	Version Control System Integration	Private module registry
Module Registry		

---

## EXERCISE

### Lab 3: Getting Started with TFC

Duration: 20 minutes

Goal: Sign up and get familiar with Terraform Cloud by creating an organization and workspace

# Remote State Data and TFC

```
terraform {
  required_version = "1.1.0"

  required_providers = {

    aws = {
      version = "~>3.0"
      source  = "hashicorp/aws"
    }
  }

  backend "s3" {
  }
}
```

```
terraform {  
  backend "s3" {  
    bucket = "mybucket"  
    key    = "path/to/my/key"  
    region = "us-east-1"  
  }  
}
```

 TERMINAL

```
> terraform init --backend-config="secret_key=ASDFGHJJUH"  
  
> export AWS_SECRET_ACCESS_KEY=ASDFGHJKIHN
```

 TERMINAL

```
> terraform init --backend-config="secret_key=ASDFGHJJUH"
```

Initializing modules...

Initializing the backend...

Do you want to copy existing state to the new backend?

Pre-existing state was found while migrating the previous "local" backend to the newly configured "remote" backend. No existing state was found in the newly configured "remote" backend. Do you want to copy this state to the new "remote" backend? Enter "yes" to copy and "no" to start with an empty state.

Enter a value: yes

---

## EXERCISE

## Lab 4: Remote State

Duration: 15 minutes

Goal: Migrate your current state to a Terraform Cloud workspace

# Variables in TFC

---

## EXERCISE

## Lab 5: Secure Variables TFC

Duration: 10 minutes

Goal: Migrate your variable values to the TFC workspace and apply the configuration

# Data Sources and Conditional Logic

```
data "aws_ami" "ubuntu_20_04" {

    most_recent = true

    filter {

        name     = "name"

        values  = ["ubuntu/images/hvm-ssd/ubuntu-focal-20.04-amd64-server-*"]

    }

    owners = ["099720109477"]

}
```

```
resource "aws_instance" "web" {  
    ami                  = data.aws_ami.ubuntu_18_04.image_id  
    instance_type        = "t2.micro"  
    subnet_id            = var.subnet_id  
    vpc_security_group_ids = var.vpc_security_group_ids
```

```
condition ? true_val : false_val
```

```
count = var.create_server ? 1 : 0
```

```
(var.server_os == "ubuntu") ? ami-1234 : ami-8765
```

---

## EXERCISE

## Lab 6: Data Conditions

Duration: 15 minutes

Goal: Add data sources to your configuration and use conditional logic

# Local Values

```
locals {  
  
    service_name = "Automation"  
  
    owner        = "Cloud Team"  
  
    common_tags = {  
  
        Service = local.service_name  
  
        Owner   = local.owner  
  
    }  
  
}
```

---

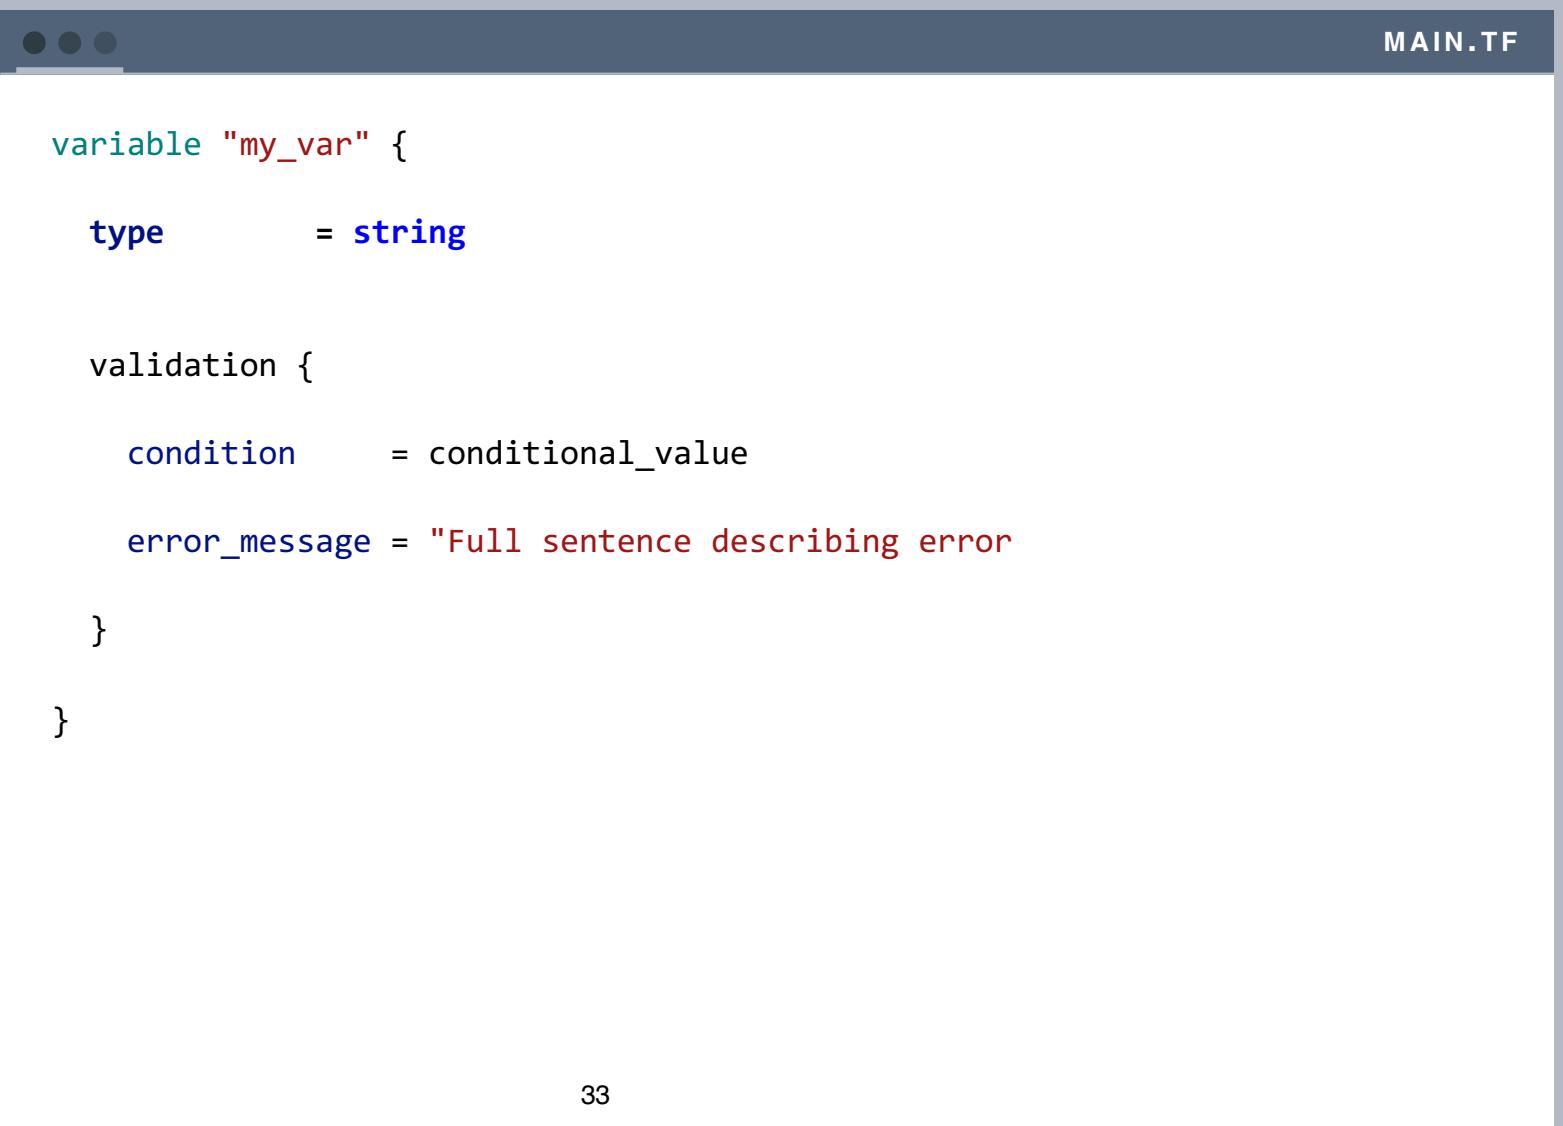
## EXERCISE

## Lab 7: Local Variables

Duration: 15 minutes

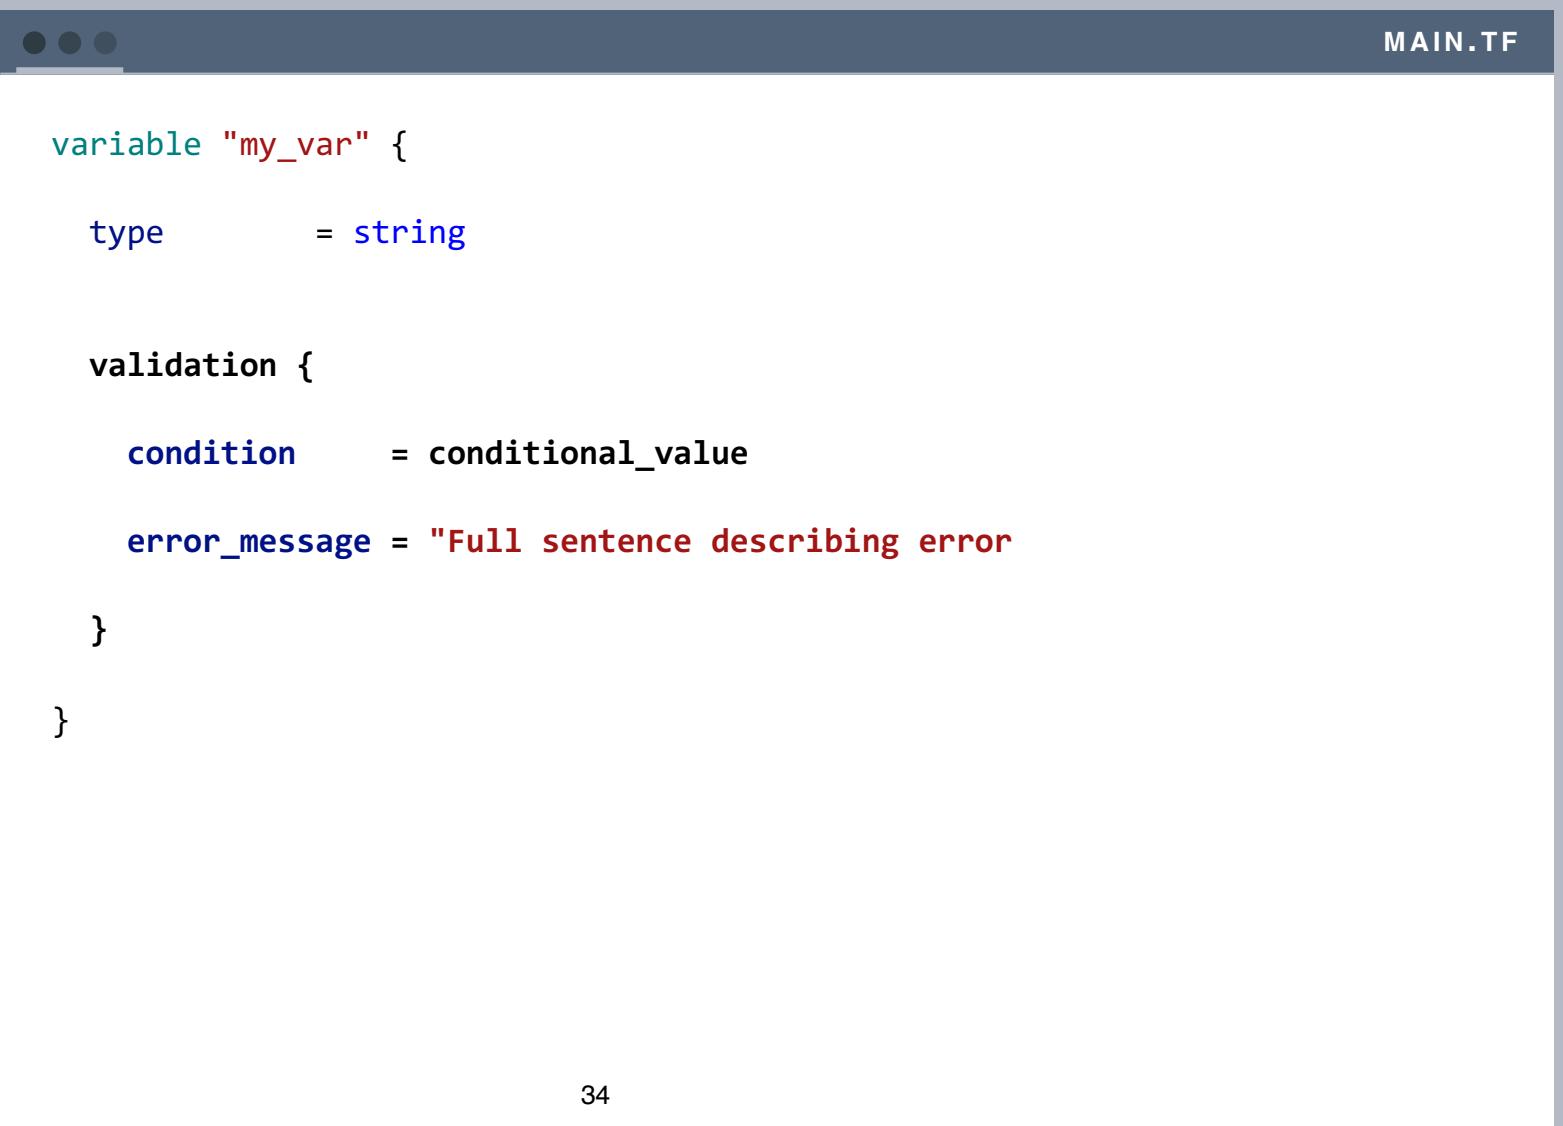
Goal: Create local variables in your configuration  
and reference them

# Variable Validation and Suppression



## MAIN.TF

```
variable "my_var" {  
    type      = string  
  
    validation {  
        condition      = conditional_value  
        error_message = "Full sentence describing error  
    }  
}
```



## MAIN.TF

```
variable "my_var" {  
    type      = string  
  
    validation {  
        condition      = conditional_value  
        error_message = "Full sentence describing error  
    }  
}
```

```
variable "server_os" {  
  
    type      = string  
  
    description = "Server Operating System"  
  
    default    = "ubuntu_20_04"  
  
  
    validation {  
  
        condition    = contains(["ubuntu_20_04", "ubuntu_18_04", "windows_2019"],  
lower(var.server_os))  
  
        error_message = "You must use an approved operating system. Options are ub  
untu_18_04, ubuntu_20_04, or windows_2019."  
  
    }  
  
}
```

```
variable "secret_key" {  
    type = string  
    sensitive = true  
}  
  
output "secret_key" {  
    value = nonsensitive(var.secret_key)  
}
```

---

## EXERCISE

## Lab 8: Variable Validation and Suppression

Duration: 20 minutes

Goal: Add validation to a variable and explore  
sensitive variables

# For-Each and Dynamic Blocks

# Looping Expressions

# Looping Constructs

Terraform HCL is a declarative language with a few primitives that allow certain types of loops, if-statements and other logic.

Terraform offers several looping constructs:

- `count` - loop over resources
- `for_each` - loop over resources and inline blocks
- `for` - loop over lists and maps

# Meta-Arguments

---

## Count

The count argument allows for N number of identical resources to be created. This removes the need for iteration with "for" or "while" loops in many cases

```
# ...

resource "aws_instance" "web" {
    count          = 2

    ami            = "ami-07669fc90e6e6cc47"
    instance_type = "t2.micro"

# ...
}
```

```
# ...  
  
output "public_ip" {  
    value = aws_instance.web.*.public_ip  
}  
  
output "public_dns" {  
    value = aws_instance.web.*.public_dns  
}
```

```
module "server" {

    count          = 2
    source         = "./server"
    ami            = ""
    server_os      = var.server_os
    subnet_id      = var.subnet_id
    vpc_security_group_ids = var.vpc_security_group_ids
    identity       = var.identity
}
```

 TERMINAL

```
> terraform state list

module.server[0].aws_instance.web
module.server[1].aws_instance.web
```

# Meta-Arguments

---

## for\_each

The `for_each` meta-argument accepts a map or a set of strings, and creates an instance for each item in that map or set.

Each instance has a distinct infrastructure object associated with it, and each is separately created, updated, or destroyed when the configuration is applied.

# The `for_each` Expression

---

```
resource "<PROVIDER>_<TYPE>" "<NAME>" {  
    for_each = <COLLECTION>  
  
    [CONFIG ...]  
  
}
```

# The each Object

---

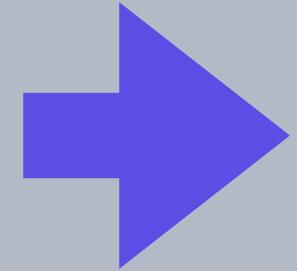
## each

Allows you to modify the configuration of each instance.

This object has two attributes:

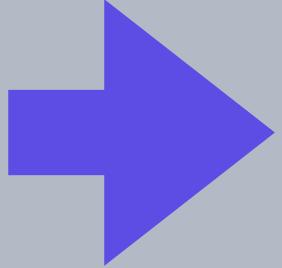
`each.key` – The map key (or set member) corresponding to this instance.

`each.value` – The map value corresponding to this instance. (If a set was provided, this is the same as `each.key`.)



```
resource "aws_iam_user" "the-accounts" {
  for_each = toset( ["Todd", "James", "Alice", "Dottie"] )
  name     = each.key
}
```

```
resource "azurerm_resource_group" "rg" {  
    for_each = {  
        a_group = "eastus"  
        another_group = "westus2"  
    }  
    name      = each.key  
    location = each.value  
}
```



```
module "server" {

    for_each                  = toset(["windows_2019","ubuntu_20_04"])

    source                    = "./server"

    ami                       = ""

    server_os                = each.value

    subnet_id                = var.subnet_id

    vpc_security_group_ids = var.vpc_security_group_ids

    identity                 = var.identity

}
```

TERMINAL

```
> terraform state list

module.server["windows_2019"].aws_instance.web
module.server["ubuntu_20_04"].aws_instance.web
```

# The `for` Expression

---

`for`

Terraform's `for` expressions also allow you to loop over a list or a map using the following syntax.

**Loop over a list / produce a list**

```
[ for <ITEM> in <LIST> : <OUTPUT> ]
```

**Loop over a map / produce a list**

```
[ for <KEY>, <VALUE> in <MAP> : <OUTPUT> ]
```

What data type  
should be  
returned?

What data source  
are we using?

```
[ for list_items in local.list_data : data_in_list ]
```

How will I refer to  
elements in the  
data source?

What are the  
elements in the  
data type  
returned?

```
[ for keys, values in local.map_data : data_in_list ]
```

*What data type  
should be  
returned?*

```
{ for key, value in local.map_data : new_key => new_values }
```

*What data source  
are we using?*

*How will I refer to  
elements in the  
data source?*

*What are the  
elements in the  
data type  
returned?*

```
{ for value in local.list_data : new_key => new_values }
```

# Looping

	Use Case	Expression
count	Multiple copies of an entire resource	<code>count = 2</code>
for_each	Multiple copies of an entire resource Multiple copies of an inline block within a resource	<code>resource "&lt;PROVIDER&gt;_&lt;TYPE&gt;" "&lt;NAME&gt;" {     for_each = &lt;COLLECTION&gt;     [CONFIG ...] }</code>
for	Generate values by looping over lists and/or maps	<code>[for &lt;ITEM&gt; in &lt;LIST&gt; : &lt;OUTPUT&gt;] [for &lt;KEY&gt;, &lt;VALUE&gt; in &lt;MAP&gt; : &lt;OUTPUT&gt;]</code>

---

## EXERCISE

## Lab 9: For Each

Duration: 20 minutes

Goal: Observe the use of a count meta-argument  
and replace it with for-each

# Dynamic Blocks

# Looping

Sometimes you need to create multiple inline blocks within a resource.

You can dynamically construct repeatable nested blocks using Terraform dynamic blocks

A dynamic block acts much like a for expression, but produces nested blocks instead of a complex typed value. Utilized the `for_each` expression.

# Multiple Inline Blocks

---

```
resource "aws_security_group" "main" {
  name      = "core-sg"
  vpc_id    = aws_vpc.vpc.id

  ingress {
    description = "Port 443"
    from_port   = 443
    to_port     = 443
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    description = "Port 80"
    from_port   = 80
    to_port     = 80
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
}
```

# Dynamic Blocks

```
dynamic "<VAR_NAME>" {  
    for_each = <COLLECTION>  
  
    content {  
        [CONFIG...]  
  
    }  
}  
}
```

# List of Maps

---

```
locals {  
    ingress_rules = [ {  
        port          = 443  
        description = "Port 443"  
    },  
    {  
        port          = 80  
        description = "Port 80"  
    }  
]
```

# Dynamic Block

---

```
resource "aws_security_group" "main" {
  name      = "core-sg"
  vpc_id    = aws_vpc.vpc.id

  dynamic "ingress" {
    for_each = local.ingress_rules

    content {
      description = ingress.value.description
      from_port   = ingress.value.port
      to_port     = ingress.value.port
      protocol    = "tcp"
      cidr_blocks = ["0.0.0.0/0"]
    }
  }
}
```

```
dynamic "block_name" {  
  
    for_each = list | map | set  
  
    iterator = iterator_name  
  
    content {  
  
        key = iterator_name.value  
  
        key = expression...  
  
    }  
  
}
```

```
resource "aws_security_group" "main" {

    name      = "core-sg"

    dynamic "ingress" {

        for_each = local.ingress_rules

        content {

            description = ingress.value.description

            from_port   = ingress.value.port

            to_port     = ingress.value.port

            protocol    = "tcp"

            cidr_blocks = ["0.0.0.0/0"]
        }
    }
}
```

---

## EXERCISE

## Lab 10: Dynamic Block

Duration: 15 minutes

Goal: Create security group rules with a dynamic block

# Private Module Registry and Version Control

---

## EXERCISE

## Lab 11: Private Module Registry

Duration: 10 minutes

Goal: Explore the private module registry in TFC  
and add a module from the public registry

---

## EXERCISE

## Lab 12: Connect VCS

Duration: 10 minutes

Goal: Connect TFC to GitHub using Oauth for  
VCS integration

---

## EXERCISE

## Lab 13: Publish Module to Registry

Duration: 20 minutes

Goal: Create a private module in GitHub and publish to the private registry

---

## EXERCISE

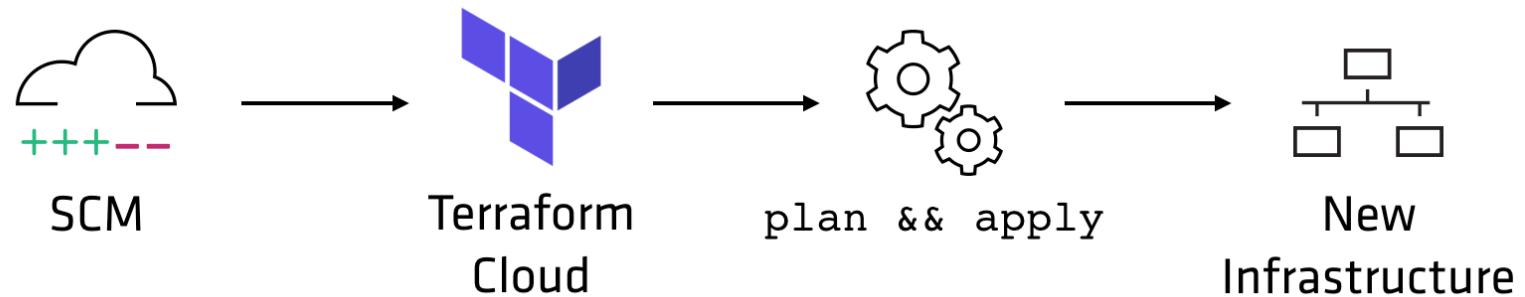
## Lab 14: Consume Registry

Duration: 10 minutes

Goal: Use the configuration designer with modules from the private registry

# TFC VCS Workflow

# VCS Workflow



---

## EXERCISE

## Lab 15: VCS Workspace

Duration: 20 minutes

Goal: Link a TFC workspace to a VCS repository

# Null Resource and Remote State Data

```
resource "null_resource" "web_app" {

    # Triggers define what changes should trigger the null resource

    triggers = { }

    # Run provisioners

    connection { }

    provisioner "file" {

    }

    provisioner "remote-exec" {

    }

}
```

```
data "terraform_remote_state" "state_name" {  
    backend = "remote"  
  
    config = {  
        hostname = "app.terraform.io"  
        organization = "org_name"  
  
        workspaces = {  
            name = "workspace_name"  
        }  
    }  
}
```

---

## EXERCISE

## Lab 16: Null Resources

Duration: 15 minutes

Goal: Use a null resource to execute scripts based on a trigger



---

Terraform 201

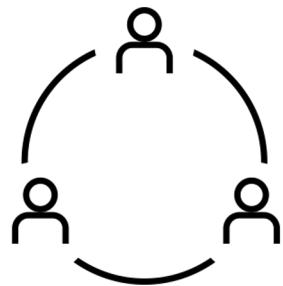
# Code Organization for Teams

# What You'll Learn

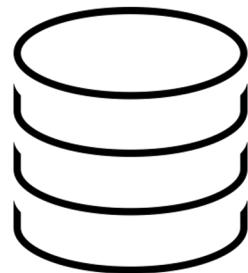


- ▢ How Terraform projects evolve
- ▢ When to use modules vs a new project
- ▢ How to use shared state to manage complexity

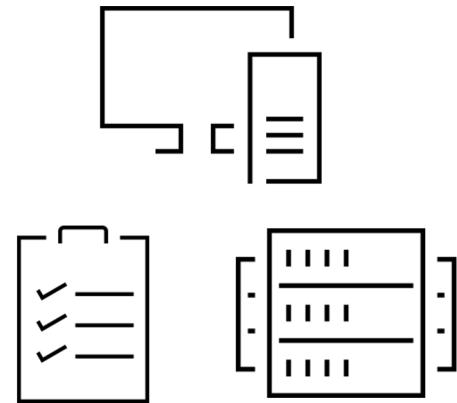
# Ways a Project Grows



Larger Team



More Resources



Dev, Test, Staging, Prod

# Discussion



- *Is this config owned by one team?*
- *Is this config used by more than one team?*
- *Does this config create resources?*

# Complexity, Capability



- A more expansive infrastructure will need a more complex Terraform project layout, but also requires more communication and additional layers of code management.
- Start simple. Refactor to manage technical or organizational complexity.



## The Growth Continuum



### Early Adoption

Dynamic

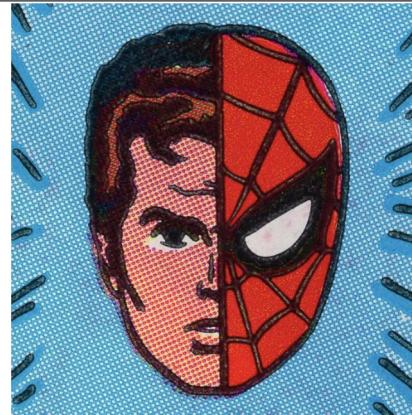
Explicit

Shared

Customizable

Creative

Simple



### Mature Use

Stable

DRY

Isolated

Reusable

Consistent

Complex

## The Growth Continuum



### Early Adoption

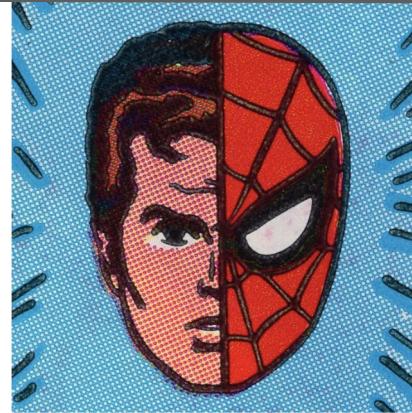
Single Local State

Root Module

Single Directory

Freeform Code

Full Control



### Mature Use

Many Remote State Files

Versioned Modules

Many Repositories

Limited Configuration

Limited Access

# Evolution of a Terraform Project



1. Single File
2. Separate Environments
3. State per Environment
4. Modules
5. Nested Modules
6. Distributed State
7. Repository Isolation



# Single file

main.tf



The image shows a dark-themed terminal window with three dots at the top left and the word "TERMINAL" at the top right. Inside the terminal, there is a file tree structure:

```
1-single-file
├── main.tf
└── terraform.tfstate
    └── terraform.tfvars
```



# Multiple Environments

May contain test and prod  
("web\_test", "web\_prod")

Possibly in different files  
within a single repo.

Lowest barrier to entry for  
testing changes.



The terminal window displays a file structure for managing multiple environments using Terraform. The directory tree is as follows:

```
2-environments
├── main-prod.tf
├── main-test.tf
└── terraform.tfstate
└── terraform.tfvars
```



## Separate States

Shrinks “blast radius” of changes between infrastructures.

Control timing and release schedule of provisioning

Provides ability to isolate resources and secrets.

```
  └── 3-states
      ├── prod
      │   ├── main.tf
      │   ├── terraform.tfstate
      │   └── terraform.tfvars
      └── test
          ├── main.tf
          ├── terraform.tfstate
          └── terraform.tfvars
```



# Resources in Modules

Improves long term maintenance

Eases reuse

Encourages consistency in design

TERMINAL

```
— 4-modules
  └── envs
      ├── prod
      │   ├── main.tf
      │   ├── terraform.tfstate
      │   └── terraform.tfvars
      └── test
          ├── main.tf
          ├── terraform.tfstate
          └── terraform.tfvars
  └── modules
      ├── core
      │   ├── input.tf
      │   ├── main.tf
      │   └── output.tf
      ├── db
      ├── network
      └── webapp
```



# Nested Modules

Groups functionality of modules into logical slices.

Allows hierarchical organization of Terraform code.

Provides the ability to effectively namespace modules.

Allows using broad or concise pieces of larger modules.

```
— 5-nested
  └─ envs
    └─ prod
      └─ main.tf
      └─ terraform.tfstate
      └─ terraform.tfvars
    └─ test
      └─ main.tf
      └─ terraform.tfstate
      └─ terraform.tfvars
  └─ modules
    └─ common
      └─ aws
        └─ compute
          └─ db
          └─ web
        └─ network
          └─ priv_subnet
          └─ pub_subnet
          └─ vpc
    └─ project
      └─ core
      └─ webapp
```



# State/Repo Per Component

Full isolation

Control team read/write permissions on each repo

Full control over access to state outputs

All apply actions are related to only the component.

```
— 6-repo-isolation
  └── envs
      ├── prod
      └── test
          ├── core
          │   └── terraform.tfstate
          ├── db
          │   └── terraform.tfstate
          ├── network
          │   └── terraform.tfstate
          └── webapp
              └── terraform.tfstate
  └── modules
      ├── common
      │   └── aws
      │       ├── compute
      │       │   ├── db
      │       │   ├── web
      │       │   └── network
      │       │       ├── priv_subnet
      │       │       ├── pub_subnet
      │       │       └── vpc
      └── project
          ├── core
          └── webapp
```

# Discussion Conclusions



- ☒, *Is this config owned by one team?*
  - ⌚ *The code should live in its own repository.*
- ☒, *Is this config used by more than one team?*
  - ⌚ *Break code down into modules where possible.*
- ☒, *Does this config create resources?*
  - ⌚ *Publish outputs to expose useful resource attributes.*



# Challenge Discussion



## Challenge



## Overview



## Questions?

### **How would you architect a project?**

Consider a larger project where you need to create compute instances, some networking, and some storage.

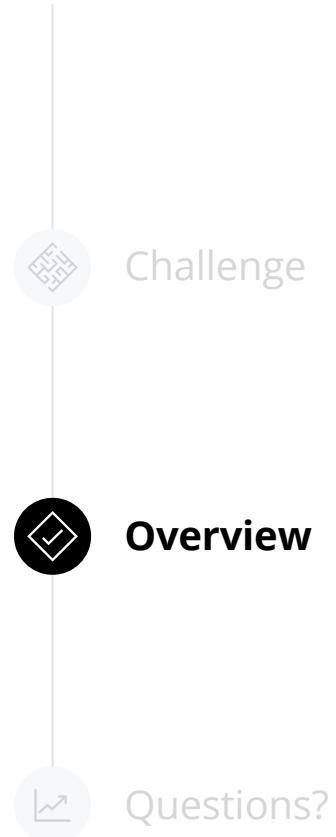
How would you divide the project into smaller configurations?

What values would need to be emitted by each configuration?

What values would need to be read by each?



# Challenge Discussion



**How Terraform projects evolve**

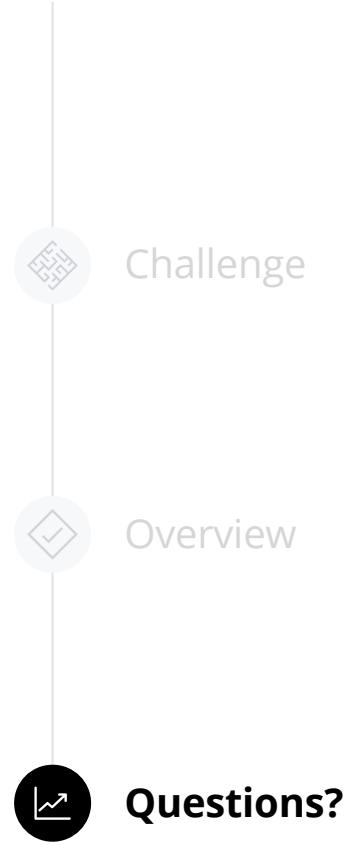
**When to use modules vs a new project**

**How to use shared state to manage complexity**



---

# Challenge Discussion



**How could you foresee these concerns impacting your environment?**



**Take a few minutes to discuss with your neighbor.**



---

Terraform 201

# Terraform Cloud Basics

# What You'll Learn



- Create a repository at GitHub
- Push your code to GitHub
- Create an organization at Terraform Cloud
- Connect to GitHub via OAuth
- Add a workspace that references a GitHub repo
- Commit changes and create infrastructure

# Terraform Cloud Features



Runs Terraform plan / apply

System of record for Terraform

Collaboration and Governance

VCS integration

Private install or hosted

<https://app.terraform.io/signup/account>

## Create an account

Have an account? [Sign in](#)

USERNAME

EMAIL

PASSWORD

 •

I agree to the [Terms of Use](#).

I acknowledge the [Privacy Policy](#).

Please review the [Terms of Use](#) and [Privacy Policy](#).

[Create account](#)

# Create a new Workspace

Workspaces allow you to organize infrastructure and collaborate on Terraform runs.

1 Connect to VCS

2 Choose a repository

3 Configure settings

## Connect to a version control provider

Choose the version control provider that hosts your Terraform source code. If you're planning to [upload your configurations via the API](#), you can [skip this step](#).



New to Terraform? [Take a tutorial](#)

## Terraform Cloud by HashiCorp would like access to:



**Your account** (robin-norwood)  
Verify your GitHub account



**Resources**  
Determine what resources both you and Terraform Cloud can access

Terraform Cloud has been installed on 1 account you have access to:  
**hashicorp**.

[Learn more about Terraform Cloud](#)

[Authorize Terraform Cloud by HashiCorp](#)

Authorizing will redirect to  
<https://app.terraform.io>

# Create a new Workspace

Workspaces allow you to organize infrastructure and collaborate on Terraform runs.

1 Connect to VCS

2 Choose a repository

3 Configure settings

## Choose a repository

Choose the repository that hosts your Terraform source code. We'll watch this for commits and pull requests.



robin-norwood 1 repository

Filter

demo-terraform-101

Can't see your repository? Enter its ID below, e.g. `acme-corp/infrastructure` :

acme-corp/infrastructure



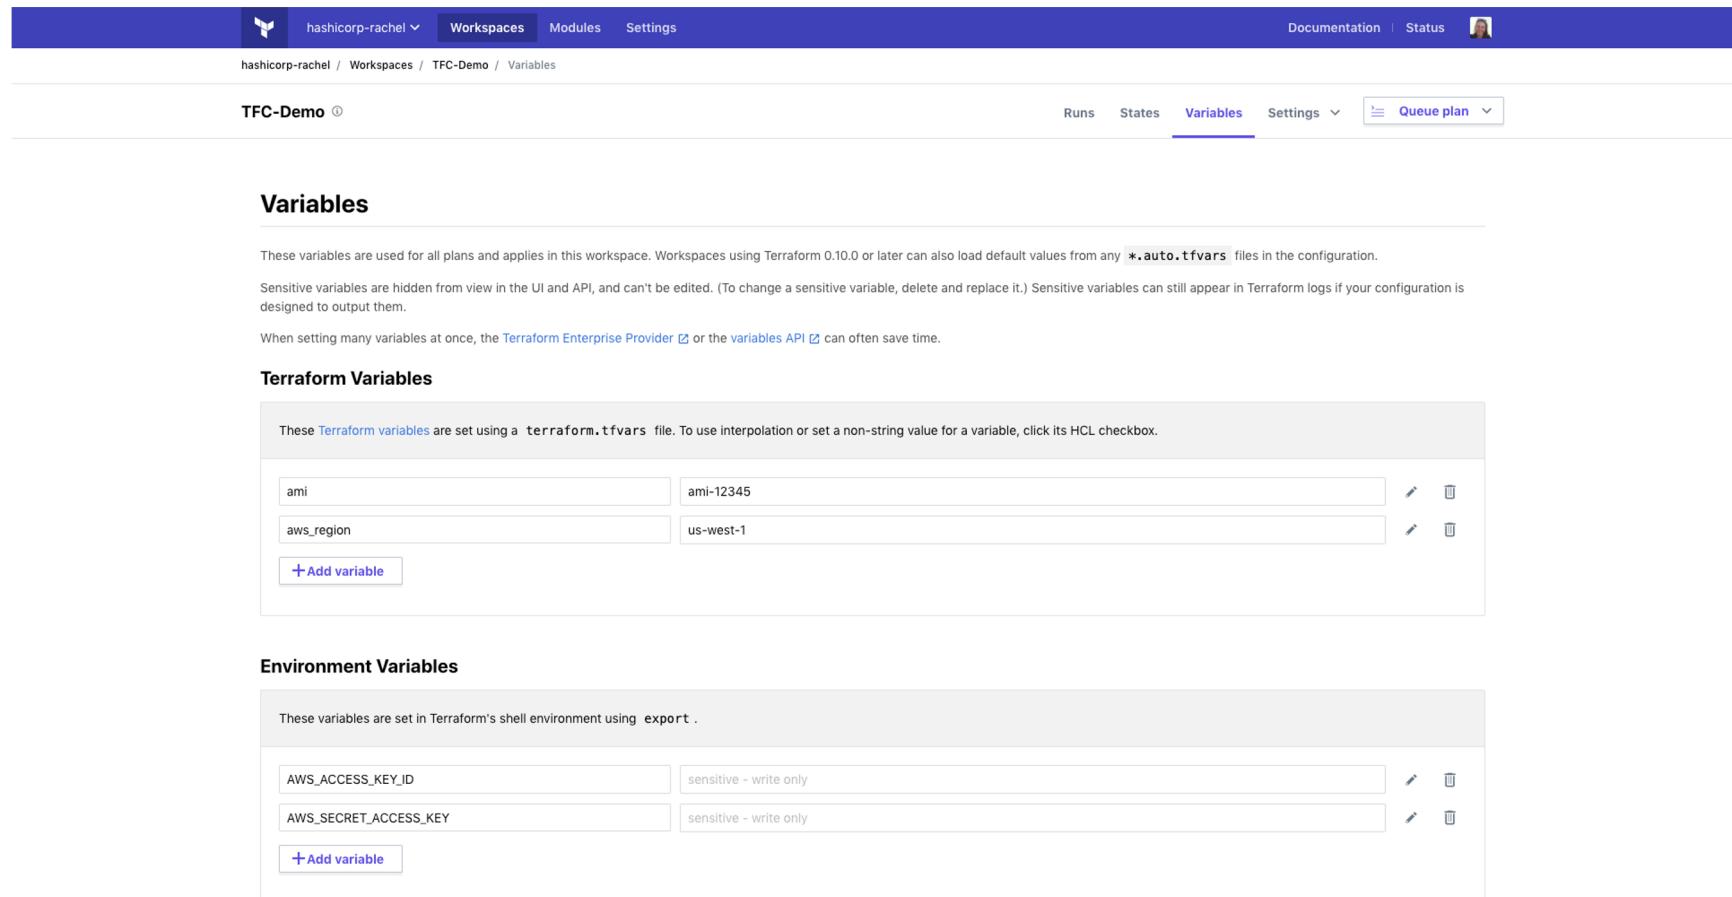


## Configuration uploaded successfully

---

Your configuration has been uploaded. Next, you probably want to configure variables (such as access keys or configuration values). If your configuration doesn't require variables, you can queue your first plan now.

[Configure variables](#)[Queue plan](#)



The screenshot shows the HashiCorp Terraform Cloud interface. At the top, there's a navigation bar with a logo, the workspace name "hashicorp-rachel", and tabs for "Workspaces", "Modules", and "Settings". On the right of the navigation bar are links for "Documentation", "Status", and a user profile icon. Below the navigation bar, the URL path is "hashicorp-rachel / Workspaces / TFC-Demo / Variables". The main content area has a header "TFC-Demo" with a dropdown arrow, and a navigation bar below it with tabs for "Runs", "States", "Variables" (which is underlined in blue), "Settings", and "Queue plan".

**Variables**

These variables are used for all plans and applies in this workspace. Workspaces using Terraform 0.10.0 or later can also load default values from any `*.auto.tfvars` files in the configuration.

Sensitive variables are hidden from view in the UI and API, and can't be edited. (To change a sensitive variable, delete and replace it.) Sensitive variables can still appear in Terraform logs if your configuration is designed to output them.

When setting many variables at once, the [Terraform Enterprise Provider](#) or the [variables API](#) can often save time.

**Terraform Variables**

These Terraform variables are set using a `terraform.tfvars` file. To use interpolation or set a non-string value for a variable, click its HCL checkbox.

ami	ami-12345		
aws_region	us-west-1		
<a href="#">+ Add variable</a>			

**Environment Variables**

These variables are set in Terraform's shell environment using `export`.

AWS_ACCESS_KEY_ID	sensitive - write only		
AWS_SECRET_ACCESS_KEY	sensitive - write only		
<a href="#">+ Add variable</a>			

hashicorp-rachel Workspaces Modules Settings Documentation Status

hashicorp-rachel / Workspaces / demo-terraform-101 / Runs

demo-terraform-101

Runs States Variables Settings Queue plan

Current Run

Update main.tf CURRENT ! NEEDS CONFIRMATION  
#run-snuLLqFRQ9MDAWU2 | tr0njavolta triggered from GitHub | Branch after-tfe | 32d06d4 a minute ago

Run List

Update main.tf CURRENT ! NEEDS CONFIRMATION  
#run-snuLLqFRQ9MDAWU2 | tr0njavolta triggered from GitHub | Branch after-tfe | 32d06d4 a minute ago

Update variables ! ERRORED  
#run-p4QBJeUptEMV5nqt | rachel triggered from Terraform Enterprise UI | Branch after-tfe | 8ab4bd0 2 minutes ago

test ! ERRORED  
#run-cuKT4LbYKiwRKM36 | rachel triggered from Terraform Enterprise UI | Branch after-tfe | 8ab4bd0 3 minutes ago

Queued manually in Terraform Enterprise ! ERRORED  
#run-8fm1JF4FdPyqGB1Y | rachel triggered from Terraform Enterprise UI | Branch after-tfe | 8ab4bd0 5 minutes ago

The screenshot shows the HashiCorp Terraform Cloud interface. At the top, there's a navigation bar with the logo, user name 'hashicorp-rachel', and links for 'Workspaces', 'Modules', and 'Settings'. On the right, there are 'Documentation' and 'Status' links, along with a user profile picture. Below the navigation, the URL 'hashicorp-rachel / Workspaces / demo-terraform-101 / Runs / run-snulLqFRQ9MDAUU2' is displayed. The main content area is titled 'demo-terraform-101' with a 'Runs' tab selected. A prominent orange box at the top says '! NEEDS CONFIRMATION' and 'Update main.tf'. Below this, a message from 'trOnjavolta' indicates a run was triggered from GitHub 2 minutes ago. The run status is 'Plan finished' (green checkmark) a minute ago. It shows a timeline from 'Queued' to 'Started' to 'Finished'. A download button for 'Sentinel mocks' is available, along with a note about using them for testing policies. The execution plan details the creation of an AWS instance, listing actions like creating an AMI, associating a public IP, and setting up security groups. At the bottom, a call-to-action button says 'Apply pending' with a note about confirming the plan. There are also 'Discard Run' and 'Add Comment' buttons.

Terraform Enterprise | topfunk... x +

https://app.terraform.io/app/topfunky-demo/terraform-intro-demo/runs/run-7JAgyiCQMwrRPZHW

APPLIED Merge pull request #2 from topfunky-demo/topfunky-demo-patch-2

topfunky-demo triggered a run from GitHub 5 minutes ago Run Details

Plan finished 5 minutes ago

Apply finished a few seconds ago

Queued 2 minutes ago > Started 2 minutes ago > Finished a few seconds ago

[View raw log](#) Top Bottom Expand Full Screen

```
module.server.aws_instance.web (remote-exec): User: ubuntu
module.server.aws_instance.web (remote-exec): Password: false
module.server.aws_instance.web (remote-exec): Private key: true
module.server.aws_instance.web (remote-exec): SSH Agent: false
module.server.aws_instance.web (remote-exec): Connected!
module.server.aws_instance.web: Creation complete after 27s (ID: i-0936d03b68874f0b6)

Apply complete! Resources: 2 added, 0 changed, 0 destroyed.

Outputs:

public_dns = [
  ec2-34-212-49-75.us-west-2.compute.amazonaws.com
]
public_ip = [
  34.212.49.75
]
```

State versions created:

topfunky-demo/terraform-intro-demo#sv-LdHR89hU5K9sEmNM (Dec 18, 2017 19:27:35 pm)  
topfunky-demo/terraform-intro-demo#sv-G8QTDAv9vuHKzDqA (Dec 18, 2017 19:28:05 pm)



---

Terraform 201

# Collaborating with Terraform Cloud

# What You'll Learn



- ☐ Understand the role of Terraform Cloud in modern continuous provisioning workflows
- ☐ Understand the components of Terraform Cloud: users, organizations, teams, workspaces



hashicorp-training ▾

Workspaces

Modules

Settings

Documentation | Status



hashicorp-training / Workspaces

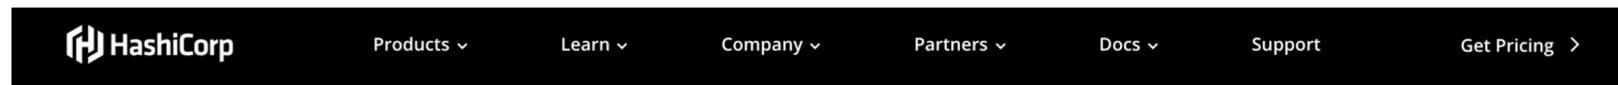
**Workspaces** 7 total[+ New workspace](#)

WORKSPACE NAME	RUN STATUS	RUN	REPO	LATEST CHANGE
consul-101-updates-testing	✓ APPLIED	run-58uH6ChRRaSQVugX	hashicorp/training	2 months ago
kait-consul-labs	! NEEDS CONFIRMATION	run-zxLMoZySs2LvNTpS	hashicorp/training	2 months ago
rachel-training			hashicorp/demo-terraform-101	6 days ago
training-lab-casey	✗ ERRORED	run-4oDEEdwTZ4A2hVRC	hashicorp/training	12 days ago
training-lab-geoffrey	✓ APPLIED	run-jgoc5XFrNKYrnjgi	hashicorp/training	6 days ago
training-lab-ryan	✗ ERRORED	run-7h7i9qjNXN5wXH6B	hashicorp/training	10 months ago
training-partner-team-prod	✗ ERRORED	run-dkFu3piezAjDWvFL	hashicorp/training-partner-team	a year ago

# What Is Terraform Cloud?



- Platform for collaboration and governance with Terraform
- Central module registry
- Connection to version control for Terraform configuration code
- Runs Terraform plan and apply
- Log of provisioning actions
- Stores state securely
- Shares state within organizations
- Configurable permissions for user and team access



# Terraform Packages

FOR INDIVIDUALS	FOR TEAMS	FOR ORGANIZATIONS
<b>Open Source</b>  Infrastructure as code provisioning and management for any infrastructure.  Free <a href="#">Download</a>	<b>Cloud</b>  Collaboration features for practitioners and small teams.  <a href="#">Sign-Up</a>	<b>Enterprise</b>  All Open Source capabilities plus collaboration and governance features for using Terraform across an organization  <a href="#">Start Trial</a>

<https://www.hashicorp.com/products/terraform/offering>



hashicorp-training ▾

Workspaces

Modules

Settings

Documentation | Status



hashicorp-training / Workspaces

**Workspaces** 7 total[+ New workspace](#)

WORKSPACE NAME	RUN STATUS	RUN	REPO	LATEST CHANGE
consul-101-updates-testing	✓ APPLIED	run-58uH6ChRRaSQVugX	hashicorp/training	2 months ago
kait-consul-labs	! NEEDS CONFIRMATION	run-zxLMoZySs2LvNTpS	hashicorp/training	2 months ago
rachel-training			hashicorp/demo-terraform-101	6 days ago
training-lab-casey	✗ ERRORED	run-4oDEEdwTZ4A2hVRC	hashicorp/training	12 days ago
training-lab-geoffrey	✓ APPLIED	run-jgoc5XFrNKYrnjgi	hashicorp/training	6 days ago
training-lab-ryan	✗ ERRORED	run-7h7i9qjNXN5wXH6B	hashicorp/training	10 months ago
training-partner-team-prod	✗ ERRORED	run-dkFu3piezAjDWvFL	hashicorp/training-partner-team	a year ago



hashicorp-rachel

Workspaces

Modules

Settings

Documentation | Status



hashicorp-rachel / Workspaces / demo-terraform-101 / Runs

demo-terraform-101 ⓘ

Runs

States

Variables

Settings

Queue plan

## Current Run

**Queued manually to destroy infrastructure**

CURRENT

✓ APPLIED

#run-9MkEA9ntc1oiFg1y | rachel triggered from Terraform Enterprise UI | Branch after-tfc | 1ffd505

a minute ago

## Run List

**Queued manually to destroy infrastructure**

CURRENT

✓ APPLIED

#run-9MkEA9ntc1oiFg1y | rachel triggered from Terraform Enterprise UI | Branch after-tfc | 1ffd505

a minute ago

**Server vars**

✓ PLANNED

#run-RW6LGKViTPRRxWMy | tr0njavolta triggered from GitHub | Branch after-tfc | 1ffd505

6 minutes ago

**Remove redundant variables**

✗ ERRORED

#run-QVGuYq9KQd25R6P9 | tr0njavolta triggered from GitHub | Branch after-tfc | 51e46f9

8 minutes ago

hashicorp-rachel Workspaces Modules Settings Documentation | Status

hashicorp-rachel / Workspaces / demo-terraform-101 / Runs / run-5cXjnLL6VweWxGpF

demo-terraform-101

Runs States Variables Settings Queue plan

✓ APPLIED Confirm variables

trOnjavolta triggered a run from GitHub 42 minutes ago Run Details

Plan finished 42 minutes ago Resources: 3 to add, 0 to change, 2 to destroy

Apply finished 40 minutes ago Resources: 3 added, 0 changed, 2 destroyed

Queued 42 minutes ago > Started 42 minutes ago > Finished 40 minutes ago

View raw log Top Bottom Expand Full screen

```
module.server.aws_instance.web[0] (remote-exec): Created symlink from /etc/systemd/system/multi-user.target.wants/webapp.service to /lib/systemd/system/webapp.service
module.server.aws_instance.web[0]: Creation complete after 1ms [id=1-0007ff1c578fd3852]

Apply complete! Resources: 3 added, 0 changed, 2 destroyed.

Outputs:

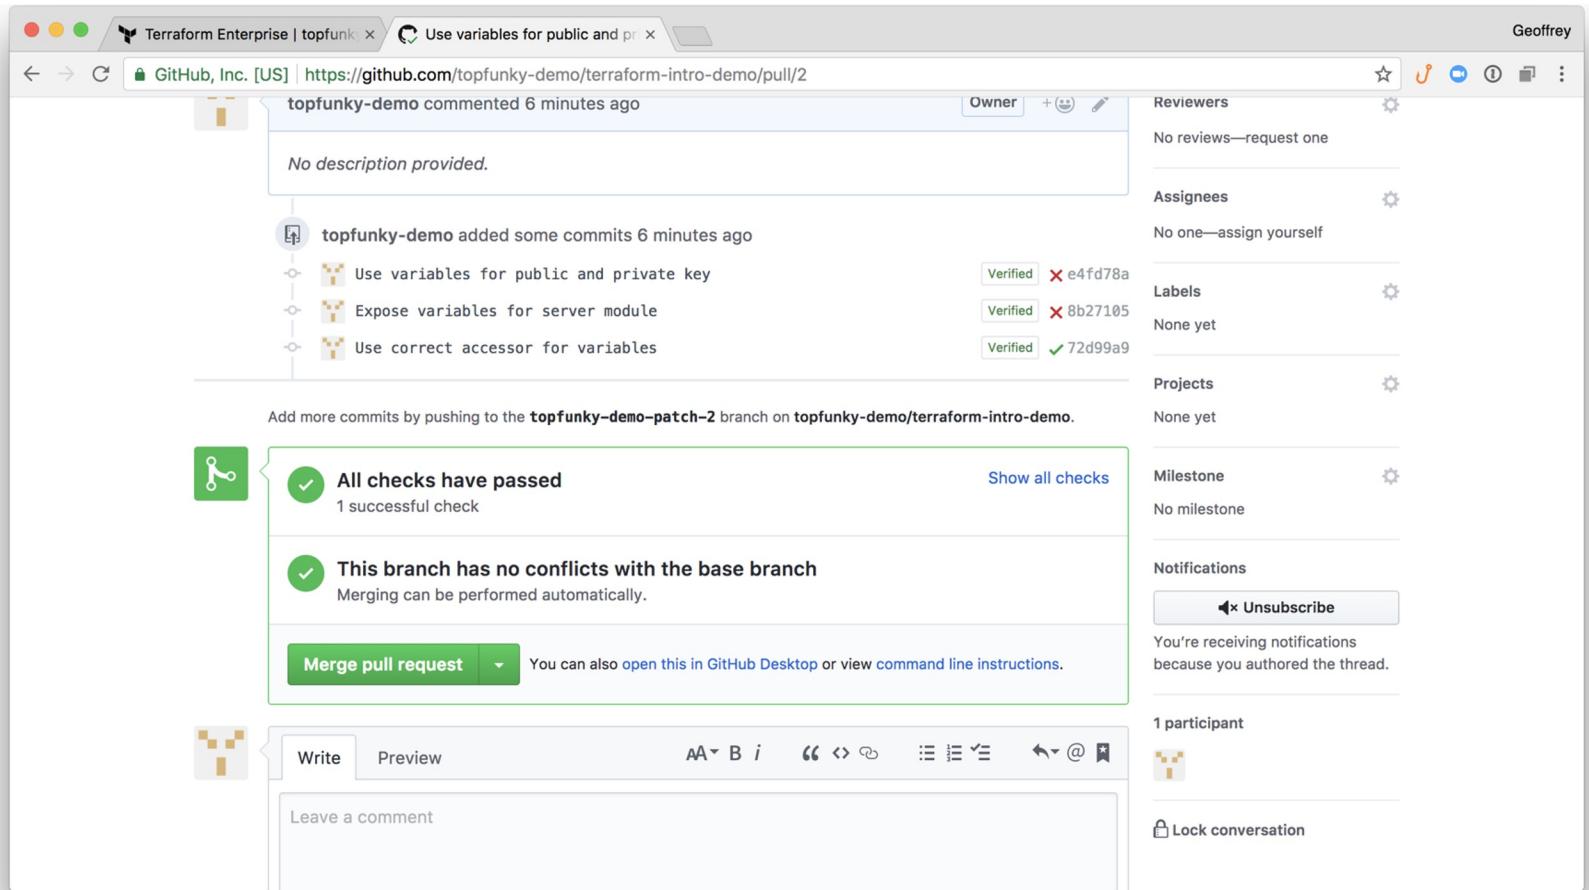
public_dns = [
  [
    "ec2-35-162-159-168.us-west-2.compute.amazonaws.com",
  ],
]
public_ip = [
  [
    "35.162.159.168",
  ],
]
```

State versions created:  
hashicorp-rachel/demo-terraform-101#sv-L5jqJTgjEtANcSx (Aug 06, 2019 22:28:36 pm)

rachel 42 minutes ago Run confirmed

Comment: Leave feedback or record a decision.

Add Comment



# User



Individual account on [app.terraform.io](https://app.terraform.io)

Can be a member of multiple teams and organizations

Signup is global at:

<https://app.terraform.io/signup>

Usernames are global within app.terraform.io.

## Create an account

Have an account? [Sign in](#)

USERNAME

EMAIL

PASSWORD

I agree to the [Terms of Use](#).

I acknowledge the [Privacy Policy](#).

Please review the [Terms of Use](#) and [Privacy Policy](#).

[Create account](#)

# Organization



Shared spaces for teams to collaborate on workspaces, policies, and Terraform modules.

Organizations can have many teams.

Sentinel policies are defined and enforced at the organization level across workspaces.



Training-Sentinel ▾

Workspaces

Settings

Documentation | Status



Training-S

hashicorp-rachel

Workspac

Training-Cloud

WORKSPACES

Training-Sentinel



Create new organization

[Get started](#)

Search by name



RUN STATUS Ø

RUN Ø

REPO Ø

LATEST CHANGE

No workspaces available yet.

# Team



Groups of users that reflect your company's organization or project structure.

Teams can be granted read, plan, write, and admin permissions to workspaces.

Teams are created under the organization settings.

Permissions are configured under workspace access settings.



hashicorp-rachel ▾

Workspaces

Modules

Settings

Documentation | Status



hashicorp-rachel / Settings / Teams

## ORGANIZATION SETTINGS

## hashicorp-rachel

General

Plan &amp; Billing

Teams

VCS Providers

API Tokens

Authentication

SSH Keys

Policies

Policy Sets

## Team Management

Teams let you group users into specific categories to enable finer grained access control policies. For example, your developers could be on a dev team that only has access to applications.

In order to allow a team access to a resource, go to the Access settings for the specific resource and enter the team name. At this point you can control the access level for that team.

The **owners** team is a special team that has implied access for all of your resources, but also has the ability to manage your organization.

## Create a New Team

NAME

Create team

## Teams

owners

3 members

# Workspaces



Workspaces are how Terraform Cloud tracks infrastructure. Workspaces contain:

- . Linked VCS repository with Terraform code.
- . Variables and values used by the configuration.
- . Current Terraform state file.
- . Historical states and run logs.

# Workspace Approaches



One approach is to have one workspace per environment per config, named accordingly.

- billing-app-dev
- billing-app-stage
- billing-app-prod
- networking-dev
- networking-stage
- networking-prod

# Alternate Workspace Approaches



Or, you could create a separate organization for development, test, or staging environments.

Organizations cannot share state, so this keeps them isolated.

This also allows modifying Terraform modules without immediately impacting production workspaces.

# Variables



Terraform Cloud allows configuring two types of variables within each workspace:

- Environment variables
- Terraform variables (contents of `terraform.tfvars` files)

Variables can be marked as *sensitive*, in which case they will not be displayed in the UI on subsequent visits.

Variables can be marked as HCL, allowing the user to define structured and typed values (maps, lists, etc.).

**training-demo** ⓘ[Runs](#) [States](#) [Variables](#) [Settings](#) ⋮ Queue plan

## Variables

These variables are used for all plans and applies in this workspace. Workspaces using Terraform 0.10.0 or later can also load default values from any `*.auto.tfvars` files in the configuration.

Sensitive variables are hidden from view in the UI and API, and can't be edited. (To change a sensitive variable, delete and replace it.) Sensitive variables can still appear in Terraform logs if your configuration is designed to output them.

When setting many variables at once, the [Terraform Enterprise Provider](#) or the [variables API](#) can often save time.

### Terraform Variables

These [Terraform variables](#) are set using a `terraform.tfvars` file. To use interpolation or set a non-string value for a variable, click its HCL checkbox.

[+ Add variable](#)

# How your team works



Terraform Cloud can be used as a part of your team's existing workflows, or you can evolve your team workflows to take advantage of Terraform Cloud features.

- **Manual:** Queue runs on demand
- **Semi-auto:** plan runs on VCS commit
- **Continuous Infrastructure:** apply on VCS commit

# API



Terraform Cloud provides a comprehensive API.  
Many companies drive workflows entirely through  
the API.

Due to the uniqueness of applications and the  
depth of features offered, we won't discuss it  
here.

See the documentation for more details:

<https://www.terraform.io/docs/cloud/api/index.html>

# API Tokens



Three kinds of API tokens are offered:  
organization, team, or user.

Accessing state from outside Terraform Cloud  
requires a user token.

Team tokens are useful for triggering plans and  
applies from external CI/CD tools.

Organization tokens only modify teams and  
workspaces. They cannot trigger plan or apply.



# Backend Configuration

```
● ● ● CODE EDITOR  
data "terraform_remote_state" "vpc" {  
  backend = "remote"  
  config {  
    organization = "company"  
  
    workspaces = {  
      name = "workspace"  
    }  
  }  
}  
# Elsewhere, use outputs from the state  
data.terraform_remote_state.vpc.outputs.cidr_block
```



---

# Discussion



**Overview**



Questions?

**Understand the role of Terraform Cloud in modern continuous provisioning workflows**

**Understand the components of Terraform Cloud: users, organizations, teams, workspaces**



---

## Discussion



Overview



Questions?



**What is the UI workflow?**



**Take a few minutes to discuss  
with your neighbor.**



---

Terraform 201

# Sentinel

# What You'll Learn



Why Sentinel?

Workflow + API

Syntax

Testing/Mocking

Data structure

Deployment

Debugging

Architecture

Common use cases

# Why Write Policy as Code with Sentinel?



Codification

Version Control

Testing

Automation

# What Sentinel Does



Policy as code

Run between plan and apply

Analyzes current state plus  
values that can be calculated  
before apply

Analyzes state of configuration

Analyzes state of system at deploy  
time

Multi-cloud, multi-provider

Augments your deployment security  
strategy

# What Sentinel Does Not Do



Continuous security checks

Audit previously deployed systems

Analyze executable contents of a VM/container/etc.

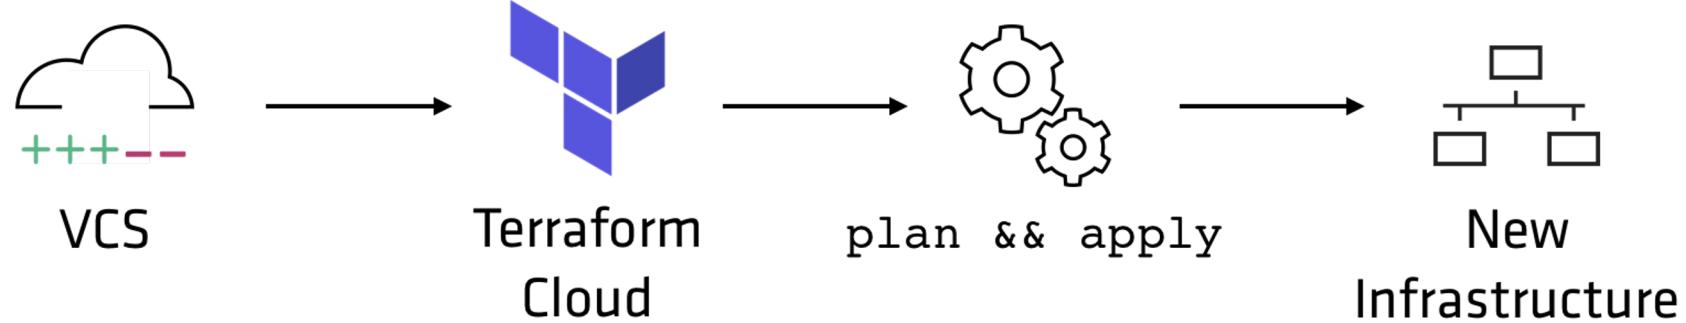
Limit runtime actions of deployed applications



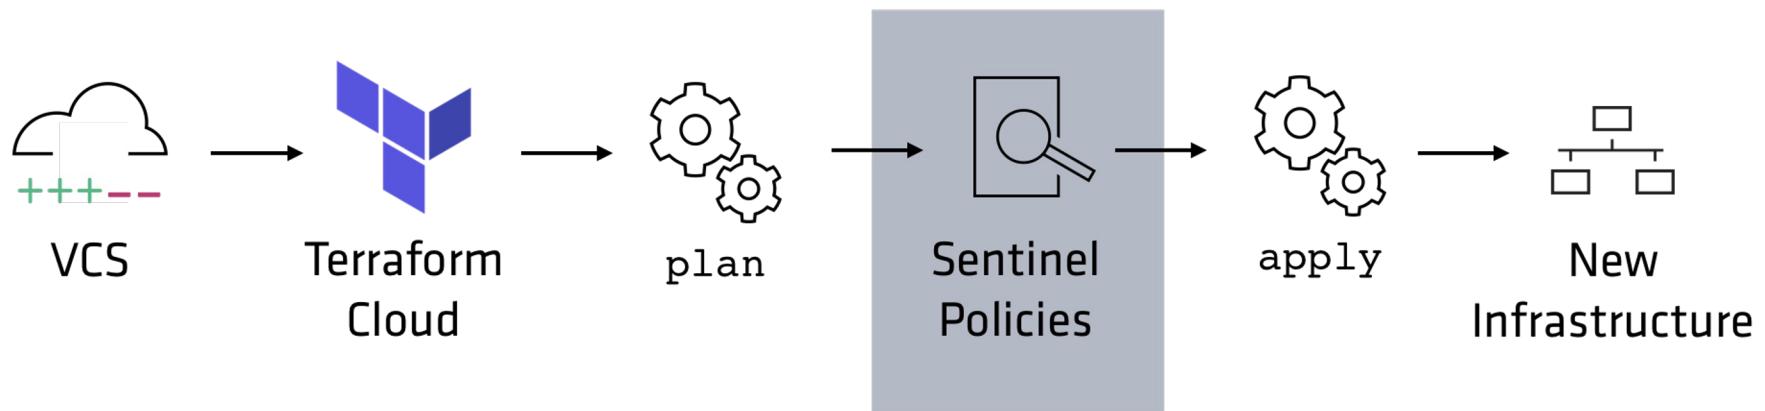
*In its current implementation, Sentinel is a tool for preventing mistakes by operators acting in good faith.*

*- Martin Atkins (Terraform Tech Lead)*

# Without Sentinel



# With Sentinel



# Become Familiar with all Sentinel Documentation



Main Sentinel docs:

<https://docs.hashicorp.com/sentinel/>

Terraform Enterprise docs:

<https://www.terraform.io/docs/cloud/sentinel/index.html>

Blog articles:

<https://www.hashicorp.com/blog/category/sentinel>

Learn Platform:

<https://learn.hashicorp.com/terraform?track=sentinel#sentinel>

# Your First Policy



## Run Task 1 in Sentinel Lab

- Download Sentinel Simulator, copy to \$PATH
- Try a simple policy named `simple.sentinel`:

```
main = rule { true }
```

A screenshot of a web browser window titled "Download Sentinel Simulator". The URL in the address bar is https://docs.hashicorp.com/sentinel/downloads. The page content includes a heading "Download Sentinel Simulator", a brief description of the tool, download links for Mac OS X (32-bit and 64-bit), and a link to FreeBSD downloads. The browser interface shows standard controls like back, forward, and search.



---

Terraform 201 - Sentinel

# Syntax

# Syntax - Introduction



- Sentinel starts with a rule named "main".
- You can define many rules with any name, but "main" is the rule which Sentinel executes first.
- The boolean return value of this rule determines the pass/fail status of the policy.

A screenshot of a dark-themed code editor window titled "CODE EDITOR". The window shows a single line of code: "main = rule { true }". The code is color-coded: "main" is pink, "rule" is purple, and the curly braces and the word "true" are white against the dark background. There are three small circular icons at the top left of the editor area, and the title bar has a standard window control icon.

# Syntax - Conditionals



- Rules can also be guarded by conditionals, which is important when writing a policy against a diverse set of infrastructure resources.
- If a conditional returns false, the rule body is not evaluated, and the result is always “true”.

A screenshot of a dark-themed code editor window titled "CODE EDITOR". The code is written in a syntax that appears to be a domain-specific language or a configuration file. It includes comments and a main rule definition.

```
# Conditional rule
# Passes if is_docker is false (rule doesn't apply)
# Otherwise evaluates has_exposed_ports
main = rule when is_docker { has_exposed_ports }
```

# Syntax - Comparators



A screenshot of a dark-themed code editor window titled "CODE EDITOR". The window contains a list of operators and their meanings:

<code>==</code>	<i>equal</i>
<code>!=</code>	<i>not equal</i>
<code>&lt;</code>	<i>less</i>
<code>&lt;=</code>	<i>less or equal</i>
<code>&gt;</code>	<i>greater</i>
<code>&gt;=</code>	<i>greater or equal</i>
<code>"is"</code>	<i>equal</i>
<code>"is not"</code>	<i>not equal</i>

# Syntax - Boolean Values



Condition:	Evaluates to:
undefined or true	true
undefined or false	undefined
undefined or undefined	undefined
undefined and true	undefined
undefined and false	undefined
undefined and undefined	undefined
undefined xor true	undefined
undefined xor false	undefined
undefined xor undefined	undefined

# Rules Are Aggressively Memoized



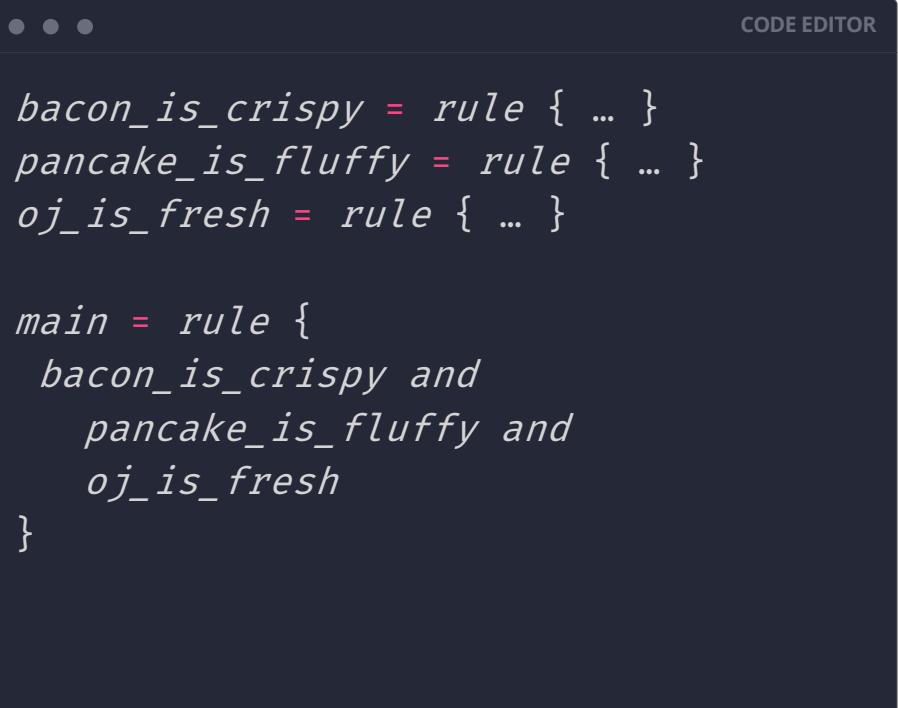
Rules are lazy and memoized (will not be recomputed unless absolutely necessary).

```
rule when Y { Z == true }
```

# Syntax Tip: Write small rules



It's slightly easier to debug policies if the rules are short and there are many of them.



A screenshot of a dark-themed code editor window titled "CODE EDITOR". The code is written in a domain-specific language, likely HCL (HashiCorp Configuration Language). It defines several small rules and a main rule that depends on them. The code is as follows:

```
bacon_is_crispy = rule { ... }
pancake_is_fluffy = rule { ... }
oj_is_fresh = rule { ... }

main = rule {
  bacon_is_crispy and
  pancake_is_fluffy and
  oj_is_fresh
}
```

# Syntax - Collections



```
all collection as index, item { } # Returns bool  
any ... # Returns bool  
for ... # NOTE: Does not return  
contains  
matches
```

# Tip: Be as Precise as You Can



Using precise comparators will determine how accurate your policies return.

For equality, use `is` rather than `contains`

# Syntax - Functions



```
a = func(x) { return x*2 }
```

Must return (or return undefined)

# Syntax - Imports



```
import "time"

main = rule { time.now.day != 1 }
```

Imports introduce additional data structures and functions for policies to leverage.

Import availability is application-specific.

Similar to imports in many programming languages.



---

Terraform 201 - Sentinel

# Testing

# Testing - Introduction



The policy itself is the test code.

JSON test file(s) may include both a fixture (mock data) and the expected result.

Multiple JSON test files may be present for each policy. Helps test policies against multiple fixtures with varying results.



A screenshot of a dark-themed code editor window titled "CODE EDITOR". The code editor displays HCL (HashiCorp Configuration Language) code. The code includes imports for "time", defines a rule for "is\_business\_hours", and sets it as the main rule. The code is as follows:

```
import "time"

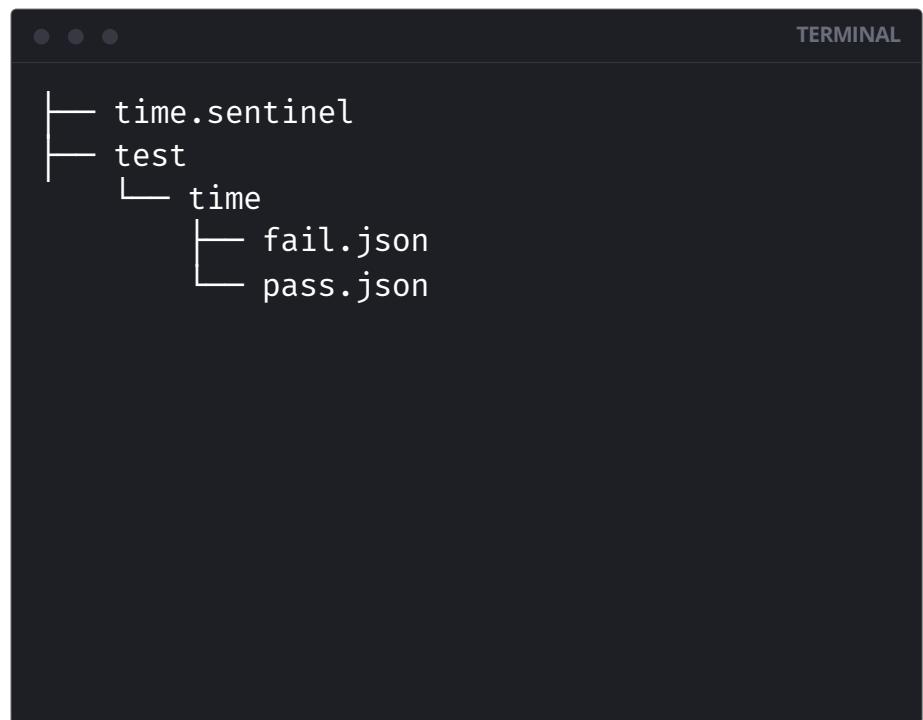
is_business_hours = rule {
    time.now.hour >= 9 and
    time.now.hour <= 17
}

main = rule { is_business_hours }
```

# Testing - Basic File Layout



- A sentinel policy file (time.sentinel)
- A test directory with a subdirectory matching the policy name (time)
- JSON files for each test case



A screenshot of a terminal window titled "TERMINAL". The window displays a file tree structure:

```
time.sentinel
test
└ time
    └ fail.json
    └ pass.json
```



## Simplest test case.

"test" defines which rules will be tested and the expected result of each.

In the previous slide, "fail.json" would expect main to return "false" here.

A screenshot of a dark-themed code editor window titled "CODE EDITOR". The editor displays a single JSON object with three properties: an opening brace {}, followed by a "test" key with a value of another object containing a "main" key with a value of "true", and finally a closing brace }.



## Other testing data

- mock: stand in for values made available by "import".



A screenshot of a dark-themed code editor window titled "CODE EDITOR". The window shows a single JSON object with several nested properties. The code is as follows:

```
{  
  "test": {  
    "main": true  
  },  
  "mock": {  
    "time": {  
      "now": {  
        "hour": 10  
      }  
    }  
  }  
}
```



## Other testing data

- global: values passed in to Sentinel from elsewhere



A screenshot of a dark-themed code editor window titled "CODE EDITOR". The window shows a single JSON configuration file. The code is as follows:

```
{  
  "test": {  
    "main": true  
  },  
  "global": {  
    "environment": "staging"  
  }  
}
```



# Sentinel Simulator

`sentinel test`  
with no arguments  
runs all test cases

```
● ● ● TERMINAL  
$ sentinel --help test  
  
Test cases are expected to be in "test/<policy>/*.json"  
files where "<policy>" is the name of the policy filename  
without an extension.  
  
$ sentinel test  
  
PASS - time.sentinel  
PASS - test/time/pass.json  PASS - test/time/fail.json
```



## Important :

### **Errors are minimal**

You might see that a test fails but not have a clear reason, even with --verbose

For example, doing import "tfplan" in a policy will cause the test to fail unless mocked (because it can't get to the import).



---

Terraform 201 - Sentinel

# Workflow in Terraform Cloud



# Sentinel Policy Sets & VCS

Managing individual policies in Sentinel is deprecated.

We now use VCS or push via the API to manage our policy sets.

The screenshot shows a GitHub repository page for 'trOnjavolta / Sentinel-Training'. The repository has 17 commits, 1 branch, 0 releases, and 1 contributor. The latest commit was made 2 hours ago. The README file contains the following text:

```
sentinel-training
An example repository for a simple Sentinel policy
```



## tags\_enforced. sentinel

This policy in the previous repo evaluates if an AWS EC2 instance is tagged and fails if not.

```
import "tfplan"

main = rule {
    all tfplan.resources.aws_instance as_,
    instances {
        all instances as _, r {
            (length(r.applied.tags) else 0) > 0
        }
    }
}
```



## sentinel.hcl

This hcl file in the same repo determines which policies to enforce and at what enforcement level.

A screenshot of a dark-themed code editor window titled "CODE EDITOR". The editor shows a single line of HCL (HashiCorp Configuration Language) code. The code defines a policy named "tags\_enforced" with an enforcement level set to "hard-mandatory".

```
policy "tags_enforced" {  
  enforcement_level = "hard-mandatory"  
}
```

# Enforcement Levels



hard-mandatory (permanent)

soft-mandatory (overridable by org owners)

advisory (log only)

A screenshot of a web browser showing the HashiCorp Settings interface for the organization "hashicorp-rachel". The browser has a standard OS X-style window title bar with red, yellow, and green buttons, a back/forward navigation bar, and a search bar.

The main header includes the HashiCorp logo, the organization name "hashicorp-rachel", and navigation links for "Workspaces", "Modules", "Settings", "Documentation", "Status", and a user profile icon.

The "Settings" menu on the left shows several categories: General, Teams, VCS Providers, API Tokens, Authentication, SSH Keys, Policies, and **Policy Sets**. The "Policy Sets" item is highlighted with a blue background and a red arrow labeled "2" pointing to it.

The main content area is titled "Policy Sets" and contains the following text:

Policy sets are groups of Sentinel policies which may be enforced on workspaces. Please see the [Sentinel in Terraform Enterprise documentation](#).

No policy sets have been created for this organization.

A prominent blue button at the top right says "Create a new policy set" with a red arrow labeled "3" pointing to it.

At the bottom of the page, there is a HashiCorp footer with links to "Support", "Terms", "Privacy", "Security", and the copyright notice "© 2019 HashiCorp, Inc.".

The screenshot shows the HashiCorp Terraform Enterprise web interface. The user is navigating through the organization settings for 'hashicorp-rachel'. They have selected the 'Policy Sets' tab, which is highlighted in blue. On the right, a modal window titled 'Create a new Policy Set' is open.

**1** In the 'Policy Set Source' section, there is a 'GitHub' button with a GitHub icon, which is highlighted with a red arrow. Below it are 'Upload via API' and a '+' button.

**2** In the 'Repository' section, a dropdown menu lists several GitHub repositories. The repository 'tr0njavolta/Sentinel-Training' is highlighted with a red arrow. Other listed repositories include 'tr0njavolta/demo-terraform-101', 'tr0njavolta/TFE-Demo', 'tr0njavolta/terraform-aws-dynamic-keys', and 'tr0njavolta/sentinel-demo'. At the bottom of the dropdown menu is a 'Create policy set' button.

At the bottom left of the interface is the HashiCorp logo. At the bottom right, there are links for 'Support', 'Terms', 'Privacy', 'Security', and a copyright notice: '© 2019 HashiCorp, Inc.'

The screenshot shows the HashiCorp Terraform Enterprise web interface. The user is navigating through the organization settings for 'hashicorp-rachel'. They are currently on the 'Policy Sets' page, which is highlighted in the sidebar.

The main content area displays a form titled 'Create a new Policy Set'. The form includes fields for 'Name' and 'Description', both of which are currently empty. Below these fields is a section for 'Policy Set Source' with three options: 'GitHub' (selected), 'Upload via API', and a '+' button.

Two red arrows are overlaid on the screenshot to indicate specific actions:

- An arrow labeled '1' points to the 'GitHub' button in the 'Policy Set Source' section.
- An arrow labeled '2' points to the repository list, specifically highlighting the entry 'tr0njavolta/Sentinel-Training'.

At the bottom of the repository list, there is a purple button labeled 'Create policy set'.

The footer of the page includes the HashiCorp logo and links for Support, Terms, Privacy, Security, and Copyright information.

**Create a new Policy Set**

**Name**  
Sentinel-Training  
You can use letters, numbers, dashes (-) and underscores (\_) in your policy set name.

**Description**

**Policy Set Source**  
[GitHub](#) [Upload via API](#) [+](#)

Create a policy set with individually managed policies

**Repository**  
tr0njavolta/Sentinel-Training  
The repository identifier in the format username/repository. Only the most recently updated repositories will appear with autocomplete; however, all repositories are available for use.

[More options \(policies path, VCS branch\)](#)

**Scope of Policies**  
 Policies enforced on all workspaces  
 Policies enforced on selected workspaces

**workspaces**

The name of the workspace you wish to add to this policy set.

demo-terraform-101

demo-terraform-101 [Remove](#)

[Add workspace](#)

**Create policy set**



## Mocked Data

Saving policies makes them immediately active for your organization's workspaces in Terraform Enterprise.

It is therefore important to test policy code locally prior to uploading.

A simple policy could be mocked with data that looks like this. This would allow the policy to be tested locally before saving it into Terraform Enterprise.

The screenshot shows a dark-themed code editor window titled "CODE EDITOR". At the top, there are three small circular icons. The main area contains the following JSON code:

```
"tfplan": {  
    "resources": {  
        "animals": {  
            "cow": [  
                { "applied": { "id":  
"animal-0" } }  
            ]  
        }  
    }  
}
```

# Examples & Documentation



HashiCorp staff frequently build and publish sample policies.

Find them in places such as the `terraform-guides` repository:

<https://github.com/hashicorp/terraform-guides/tree/master/governance/second-generation>

Read, understand, and refer to the documentation on imports (`tfplan`, `tfconfig`, `tfstate`, `tfrun`):

<https://www.terraform.io/docs/cloud/sentinel/import/index.html>



---

Terraform 201 - Sentinel

# Mock Data in Terraform Cloud

demo-terraform-101 ⓘ

Runs States Variables Settings Queue plan CURRENT

**I NEEDS CONFIRMATION Test**

rachel triggered a run from Terraform Enterprise UI a few seconds ago Run Details

**Plan finished** a minute ago Resources: 4 to add, 0 to change, 0 to destroy

Queued a few seconds ago > Started a few seconds ago

[Download Sentinel mocks](#) [View raw](#) [Sentinel mocks can be used for testing your Sentinel policies](#)

Following log Top Bottom Expand Full screen

```
+ security_groups = []
+ self            = false
+ to_port         = 80
},
]
+
name          = (known after apply)
name_prefix    = "rachel-demo"
owner_id       = (known after apply)
revoke_rules_on_delete = false
tags           = {
  "Created-by" = "Terraform"
  "Identity"   = "rachel-demo"
}
vpc_id         = (known after apply)
```

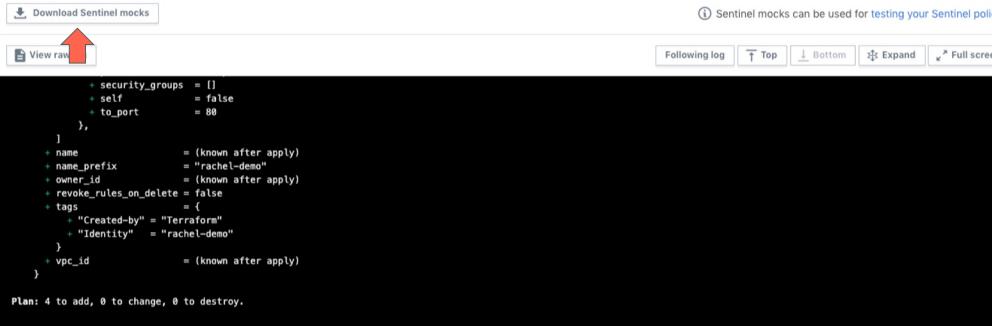
Plan: 4 to add, 0 to change, 0 to destroy.

**Policy check passed** a minute ago Policies: 1 passed, 0 failed

**Apply pending**

Needs Confirmation: Check the plan and confirm to apply it, or discard the run.

[Confirm & Apply](#) [Discard Run](#) [Add Comment](#)





# Mock data structure

```
• • • TERMINAL
sentinel-demo
├── with-data.sentinel
├── sentinel.json
└── test
    ├── fail.json
    └── pass.json
└── testdata
    ├── mock-tfconfig.sentinel
    ├── mock-tfplan.sentinel
    └── mock-tfstate.sentinel
```



# Mock data



A terminal window titled "TERMINAL" showing the contents of a JSON file named sentinel.json. The file contains a single object with a "mock" key, which points to another object containing four key-value pairs: "tfconfig", "tfplan", "tfrun", and "tfstate", each pointing to a corresponding sentinel file in the testdata directory.

```
$ cat sentinel.json

{
  "mock": {
    "tfconfig": "testdata/mock-tfconfig.sentinel",
    "tfplan": "testdat/mock-tfplan.sentinel",
    "tfrun": "testdata/mock-tfrun.sentinel",
    "tfstate": "testdata/mock-tfstate.sentinel"
  }
}

$ sentinel apply -trace instance_type_is_medium.sentinel
-config=sentinel.json`
```



---

Terraform 201 - Sentinel

# Best Practices (and a few warnings)

# Local is Better



Work on the local machine, write tests. Work in small increments if possible.

Run a syntax check with `fmt`, leverage mock data where possible.

The cycle of write-execute-observe is long in Terraform Cloud, especially if your core Terraform plan is large.

# Warnings are minimal



You won't be warned for certain types of mistakes such as scoping errors.

Writing good tests and executing them prior to a production is imperative.

# Computed values can be loopholes



Certain values may not be known until apply time.

These values cannot be effectively used in policies.

Terraform config authors could exploit this to bypass policies that would otherwise fail.

The most defensive thing you can do is disallow computed values in the policy.

# **Code defensively, favor denial as much as possible**



When possible, use a whitelist instead of a blacklist.

Example: no cidr 0.0.0.0/0 vs 192.168.0.0/16

Whitelists are more secure than blacklists, in general.

Example: 000.000.000.000/0 is a valid IP that could evade checks for 0.0.0.0/0

# The applied data structure is most useful



When examining the `tfplan` data structure, look for `applied` to find the values that will be deployed



---

Terraform 201

# Authoring with the Registry

# What You'll Learn



- Use standard module structure to build a module
- Publish a module to the registry
- Use the published module from a Terraform configuration
- Understand how to version and update a module
- Use the module configuration designer



## Modules

Collaboration, Reuse,  
Consistency

### EXAMPLES

AWS Networking  
App Server  
DB Cluster



## VCS

Access Control,  
Pull Request Workflow, CI

### EXAMPLES

GitHub  
Bitbucket (private)



## Terraform Registry

Versioning, Documentation,  
Search

### EXAMPLES

terraform-aws-networking  
terraform-go-server  
terraform-pg-cluster

# Private Module Registry



Both the hosted and private instances of Terraform include your own private module registry which can be used to securely and privately publish modules within your company.



hashicorp-rachel / Modules

## Modules

+ Design configuration

+ Add module

consul



Providers ▾

dynamic-keys PRIVATE

Terraform module that dynamically generates a public/private keypair.

Details

AWS

Version 1.0.0



## Module Registry Workflow: Git

Push code to GitHub or Bitbucket.  
The repository must be named in  
the format:

terraform-PROVIDER-NAME

Examples:

terraform-aws-peering

terraform-azure-autoscaling

The screenshot shows a GitHub repository page for 'topfunky/terraform-demo-animal'. The repository is described as a 'Sample repo for experimenting with the HashiCorp Terraform Module Registry'. It has 3 commits, 1 branch, 2 releases, and 1 contributor. The files listed are .gitignore, README.md, main.tf, outputs.tf, variables.tf, and README.md. The latest commit is from Jan 22, 2018.

File	Description	Time
.gitignore	First import	20 minutes ago
README.md	First import	20 minutes ago
main.tf	First import	20 minutes ago
outputs.tf	Rename to standard module structure	14 minutes ago
variables.tf	Sample variable. Required.	9 minutes ago
README.md		

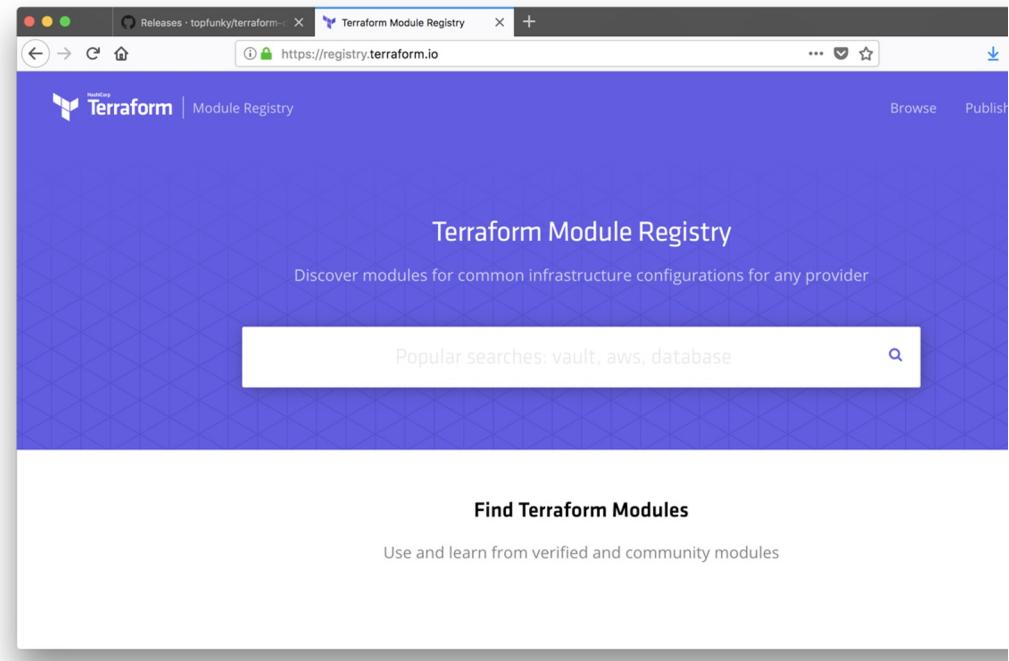


## Module Registry Workflow: Registry

Connect your SCM account to the Terraform Registry. Publish a module.

Tagged commits will be synced automatically.

v1.0.6



# Module Registry Workflow: Config



Use the code snippet from the registry to reference the module from your code. Provide all necessary variables.

Run `init` to pull in the correct version of the module.

```
terraform init
```

# Example of Modules for Registry



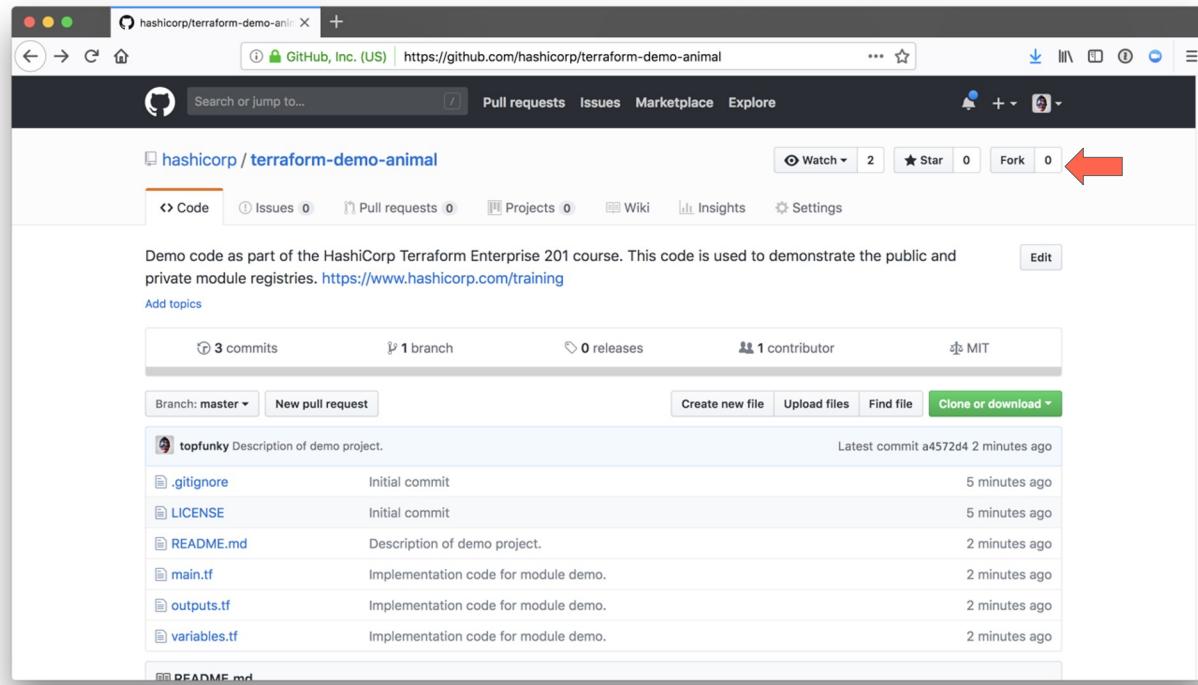
TERMINAL

```
~/projects/terraform-demo-animal
.
├── README.md
└── LICENSE
```

CODE EDITOR

```
# main.tf
resource "random_pet" "animal" {}

# outputs.tf
output "animal" {
  value      = random_pet.animal.id
  description = "Contains the name
of a random animal."
}
```



## Push tags from your Fork



```
TERMINAL

$ git clone https://github.com/$USER/terraform-
demo-animal.git

$ cd terraform-demo-animal

$ git tag v0.0.1

$ git push --tags
```

/

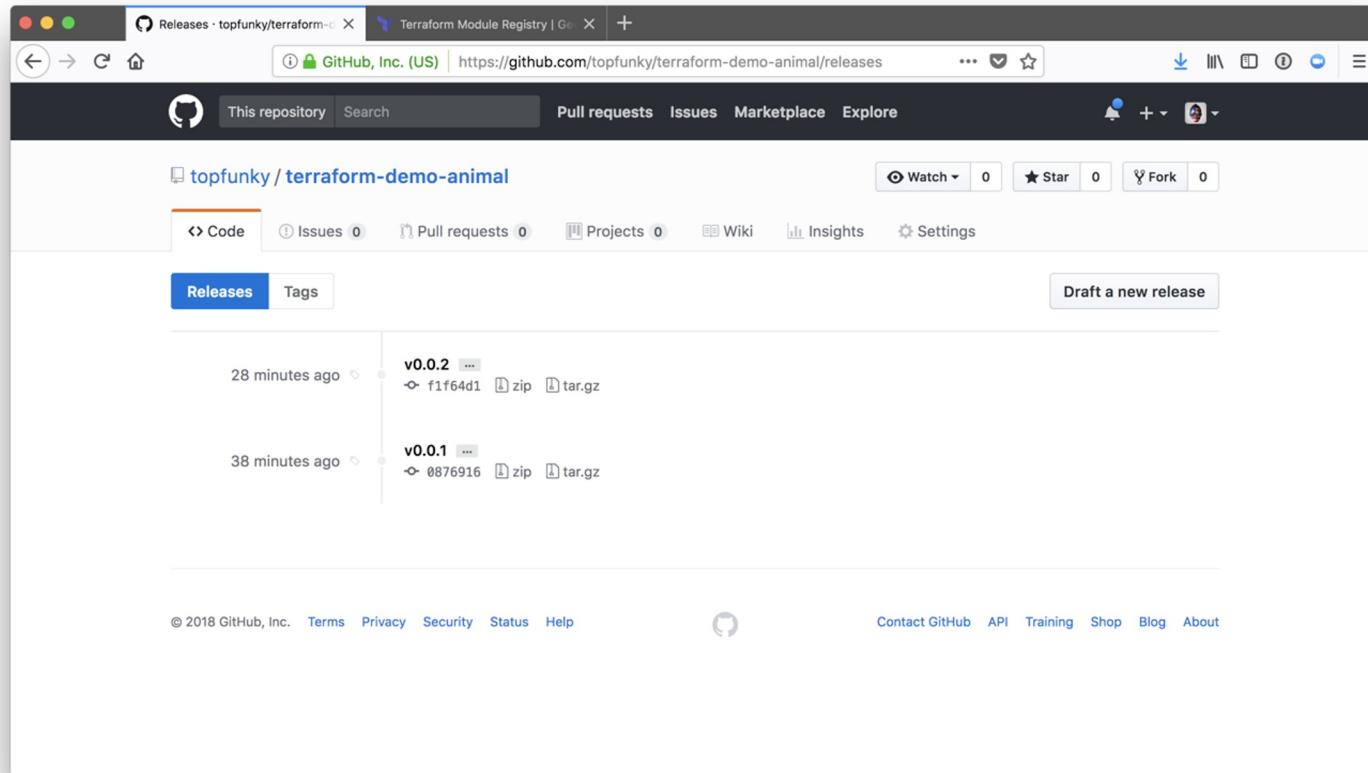
# Publish the Code on Your Private Module Registry



Go to <https://app.terraform.io/> and go to the workspaces dashboard.

Click the "Modules" button in the menu bar.

Use the "+ Add Module" button to select your repository and add it to the module registry.





---

Terraform 201 - Module Registry

# Best Practices

# Module Versioning



In your published module, don't forget to set a version for the resources and providers that you rely on.

Be sure to provide as much version flexibility as possible to make your module easy to use among various Terraform provider versions. I.e., prefer “`>= 1.0`” to “`1.0`”.

A screenshot of a dark-themed code editor window titled "CODE EDITOR". The window shows a snippet of Terraform configuration code. At the top, there are three small circular navigation dots. The code itself starts with a comment "# main.tf" and defines a provider "random" block with a "version" constraint of "~> 1.1".

```
# main.tf
provider "random" {
  version = "~> 1.1"
}
```



## Module Storage Structure

Look at the hidden ".terraform" directory and you'll see that all versions of the module are stored locally.

Metadata is in  
.terraform/modules/modules.json

```
TERMINAL
$ tree .terraform
.terraform
└── modules
    ├── 495f1ee96941ccf0a6c619930ea1160d
    │   └── topfunky-terraform-demo-animal-f1f64d1
    │       ├── README.md
    │       ├── main.tf
    │       ├── outputs.tf
    │       └── variables.tf
    ├── 753525cd62578ec39462ba136cf9f7e
    │   └── topfunky-terraform-demo-animal-0876916
    │       ├── README.md
    │       ├── main.tf
    │       └── output.tf
    └── modules.json
plugins
└── darwin_amd64
    └── lock.json
        └── terraform-provider-random_v1.1.0_x4
```



## Versioning & Dependency Metadata

The contents of .terraform/modules/modules.json show how Terraform tracks version numbers and maps them to local directories.

If you need to reset all dependencies, you can delete this directory and run "terraform init" again.

```
{  
  "Modules": [  
    {  
      "Root": "topfunky-terraform-demo-animal-0876916",  
      "Version": "0.0.1",  
      "Source": "topfunky/animal/demo",  
      "Key": "1.animal;topfunky/animal/demo.0.0.1",  
      "Dir": ".terraform/modules/753525cdd62578ec39462ba136cf9f7e"  
    },  
    {  
      "Dir": ".terraform/modules/495f1ee96941ccf0a6c619930ea1160d",  
      "Source": "topfunky/animal/demo",  
      "Key": "1.animal;topfunky/animal/demo.0.0.2",  
      "Version": "0.0.2",  
      "Root": "topfunky-terraform-demo-animal-f1f64d1"  
    }  
  ]  
}
```

# By Default, Paths are Relative to the Terraform CLI



If you want to reference resources in a published module's own directory, use

```
 ${path.module}
```



---

Terraform 201 - Module Registry

# Configuration Designer

# Configuration Designer



User interface for selecting modules and specifying their versions and required values.

Generates boilerplate Terraform configuration code to import the desired modules and pass in the appropriate values.



hashicorp-rachel ▾

Workspaces

Modules

Settings

Documentation | Status



hashicorp-rachel / Modules

## Modules



+ Design configuration

+ Add module

consul



Providers ▾

dynamic-keys

PRIVATE

Terraform module that dynamically generates a public/private keypair.

Details

AWS

Version 1.0.0

A screenshot of a Firefox browser window displaying the Terraform Enterprise Modules interface. The URL in the address bar is <https://app.terraform.io/app/modules/geoffrey-org/design/modules>. The page title is "Terraform Enterprise | Modules". The navigation bar includes links for "geoffrey-org", "Workspaces", and "Modules". On the right, there are links for "Documentation" and "Status".

The main content area shows a search bar with the text "consul" and a "Providers" dropdown menu. Below the search bar, there are two sections: "ADD MODULES TO WORKSPACE" and "SELECTED MODULES".

In the "ADD MODULES TO WORKSPACE" section, a card for the "animal" module is displayed. The card includes the following details:

- animal** (PRIVATE)
- Sample repo for experimenting with the HashiCorp Terraform Module Registry
- Details ↗
- Add Module (highlighted with a red arrow labeled "1")
- DEMO Version 0.0.2

In the "SELECTED MODULES" section, there is a green button labeled "Next »" with a red arrow labeled "2" pointing to it.

At the bottom of the page, the HashiCorp logo is on the left, and a footer navigation bar is on the right with links for Support, Terms, Privacy, Security, and the copyright notice "© 2018 HashiCorp, Inc."

Firefox File Edit View History Bookmarks Tools Window Help 100% DV Fri 2:59 PM

Terraform Enterprise | Modules X +

https://app.terraform.io/app/modules/geoffrey-org/design/variables

geoffrey-org Workspaces Modules Documentation Status

## Modules / Configuration Designer

Select Modules > Set Variables > Verify > Publish Next »

### SELECT MODULE TO CONFIGURE

animal PRIVATE

Sample repo for experimenting with the HashiCorp Terraform Module Registry

Details ↗ Configured ✓

### CONFIGURE VARIABLES

**name** REQUIRED

a name

server  Deferred

**HashiCorp** Support Terms Privacy Security © 2018 HashiCorp, Inc.

The screenshot shows a Firefox browser window with the Terraform Enterprise Modules interface. The URL in the address bar is <https://app.terraform.io/app/modules/geoffrey-org/design/verify>. The page title is "Terraform Enterprise | Modules". The top navigation bar includes links for "geoffrey-org", "Workspaces", "Modules", "Documentation", and "Status". Below the navigation, the breadcrumb trail reads "Modules / Configuration Designer". A green "Next »" button is visible on the right. The main content area is titled "Terraform Configuration" and contains the following Terraform code:

```
1 //-----
2 // Modules
3 module "animal" {
4   source  = "app.terraform.io/geoffrey-org/animal/demo"
5   version = "0.0.2"
6
7   name = "server"
8 }
```

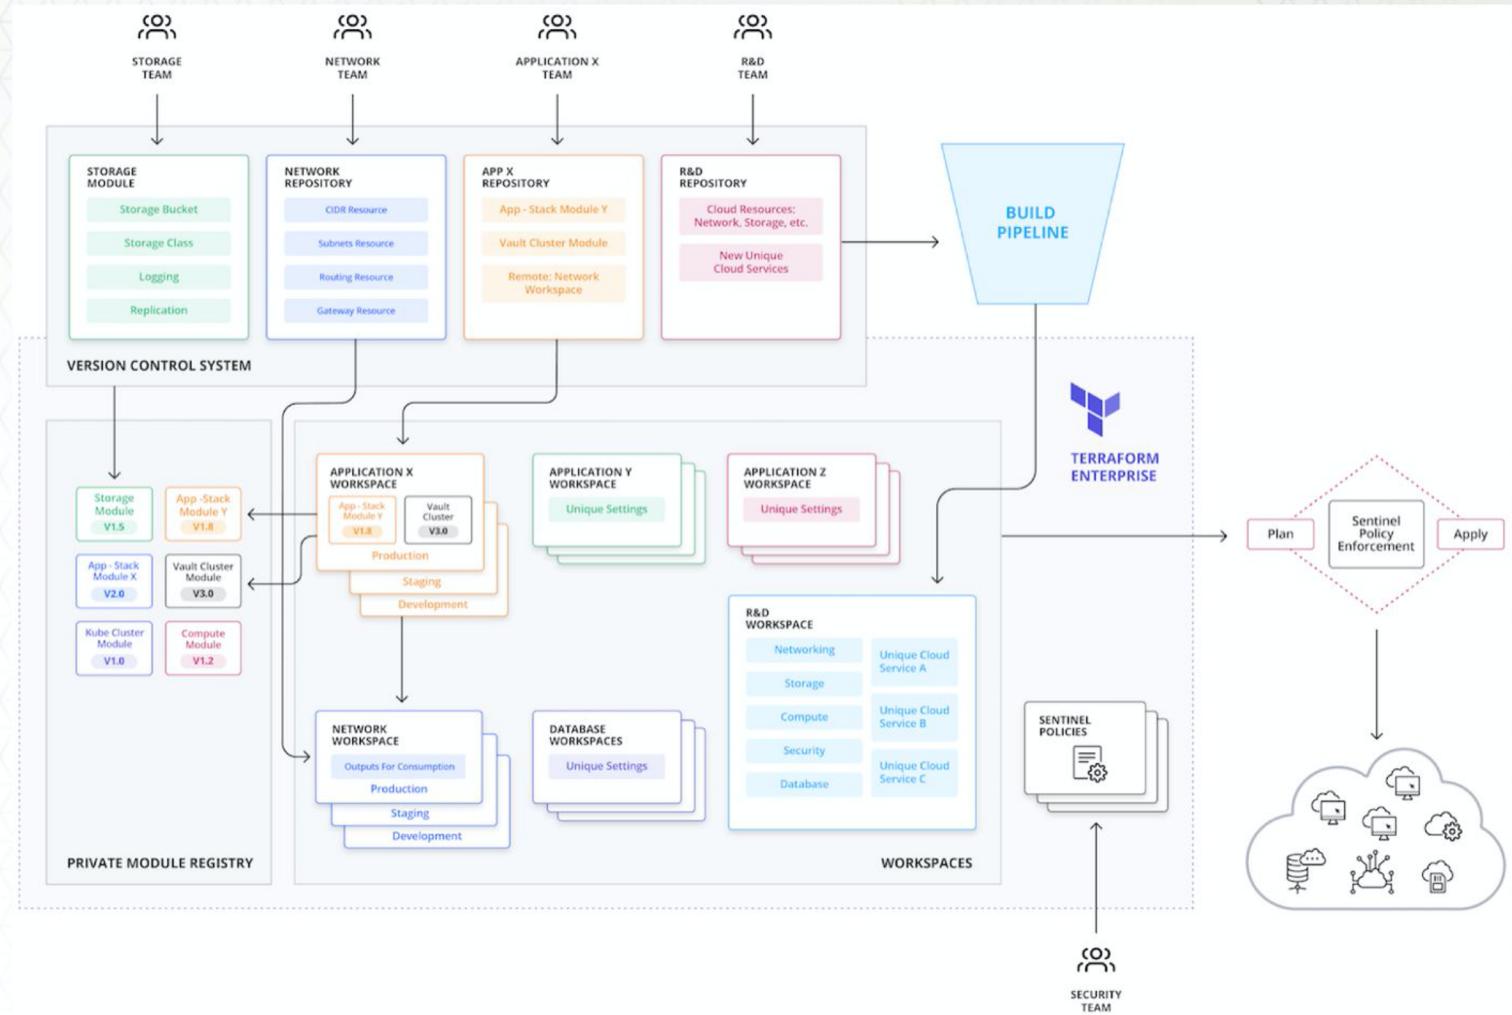
# Configuration Designer



Supports the producer/consumer pattern.

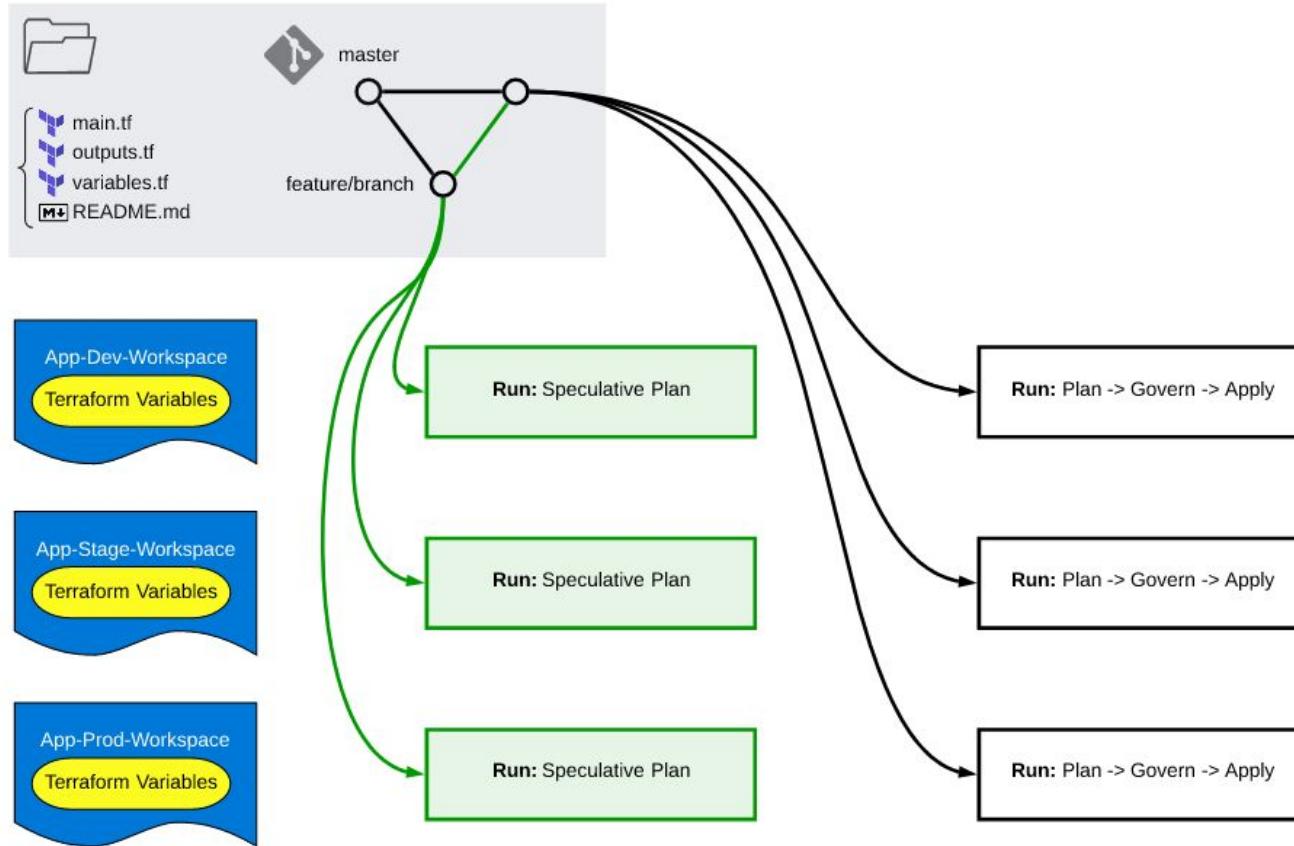
Teams can create well-documented modules with clear arguments that can be configured and used by other teams. Only minimal coding ability is required.

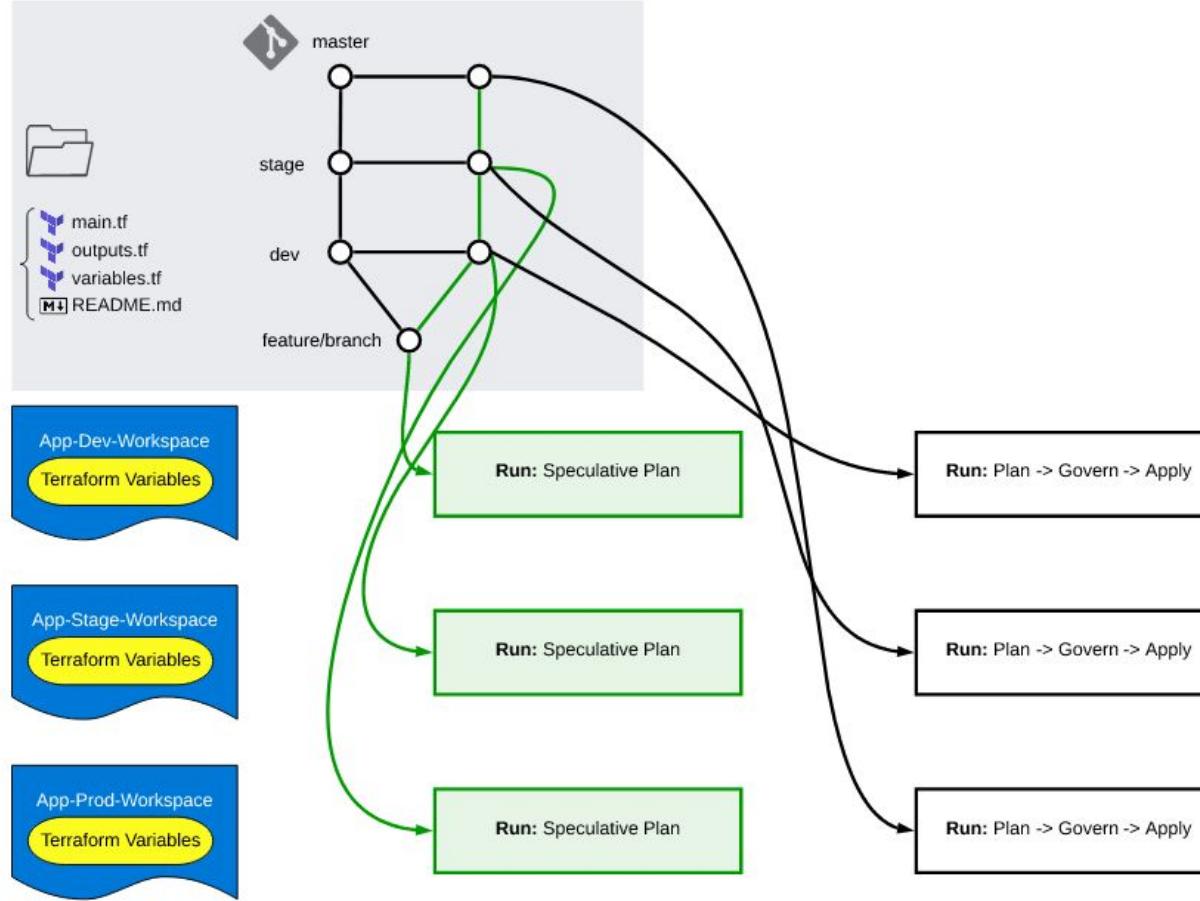
Producer teams can define best practices in code. Consuming teams don't need to understand the details, but can comply with infrastructure provisioning and security policies.

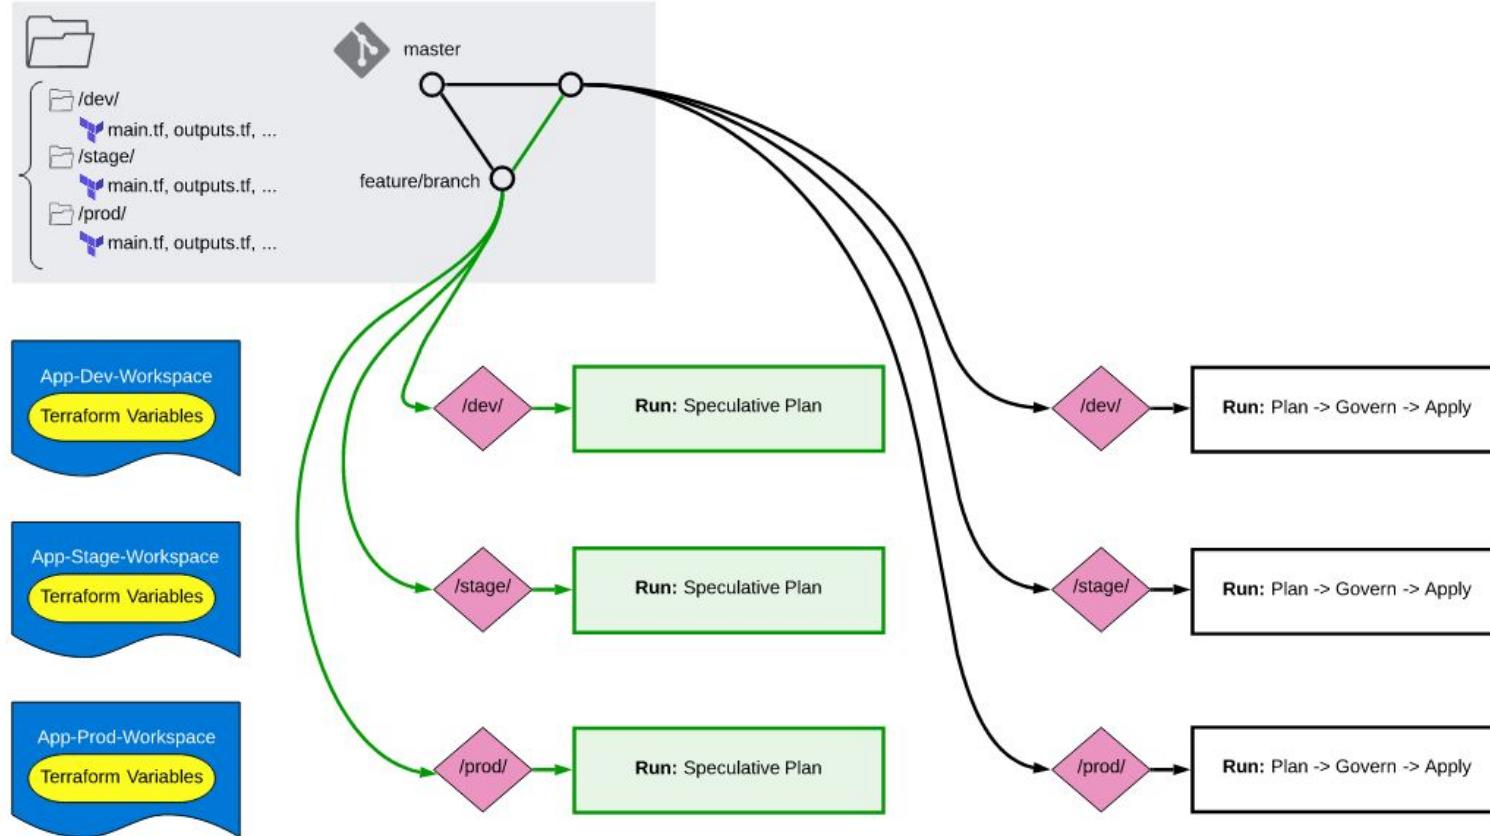


# Workspace Repo Structure Options

 HashiCorp  
**Terraform Enterprise**







# 1 Terraform Coding Best Practices

## 1.1 Purpose

This document serves as a guidepost for effective development with Terraform across multiple team members and workstreams. Our goal is to provide a unified and consistent approach to managing code by outlining some rules of engagement. These guidelines should be referenced during code reviews in order to align new development with existing work and to update old code. As a disclaimer, the guidelines found here are not law but rather best practices that we have encountered in our numerous infrastructure-as-code experiences. In cases where new needs require deviation from this guide, we should adjust the guide.

## 1.2 Intended Audience

These standards are targeted towards developers and administrators who will be writing Terraform code on a regular basis.

## 1.3 Terraform Modules

The following guidelines apply to reusable Terraform modules and root configs.

### 1.3.1 Module Structure

1. Terraform modules must follow the standard module structure.
2. Every module must start with a `main.tf` file, where resources are located by default.
  - Among other resources, all locals should be defined here in a single `locals` block.
3. All modules must have a `README.md` (with basic documentation in Markdown format).
4. Examples should be located in `examples/`, each with its own subdirectory and a `README.md` file
5. Logical groupings of resources can be grouped into their own files and given descriptive names, like `network.tf`, `instances.tf`, or `loadbalancer.tf`.
  - Avoid giving every resource its own file. Resources should be grouped by shared purpose. For example, “`google_dns_managed_zone`” and “`google_dns_record_set`” would likely be combined in `dns.tf`.
6. Only Terraform (`*.tf`) and repo metadata files (like `README.md`, `CHANGELOG.md` or `kitchen.yml`) should exist at the root directory of a module.

7. Additional documentation should be stored in a docs / subdirectory.

### 1.3.2 Naming Convention

All configuration objects should be named using underscores to delimit multiple words. This practice ensures consistency with the naming convention for resource types, data source types, and other predefined values. Note that this convention does not apply to name arguments.

```
# Good
resource "google_compute_instance" "web_server" {
    name = "web-server"
# ... }

# Bad
resource "google_compute_instance" "web-server" {
    name = "web-server"
#... }
```

### 1.3.3 Variables

1. All variables shall be declared in variables.tf.
2. Variables shall have descriptive names relevant to their usage or purpose.
  - Inputs, local variables, and outputs representing numeric values such as disk sizes or RAM size **SHOULD** be named with units (like ram\_size\_gb). APIs typically do not have standard units, so naming variables with units makes the expected input unit clear for configuration maintainers.
  - For units of storage, the unit prefix (kilo, mega, giga) **SHOULD** be binary (powers of 1024). For all other units of measurement, the unit prefix **SHOULD** be decimal (powers of 1000).
  - Boolean variables **SHOULD** be named with positive values (like enable\_external\_access) to simplify conditional logic.
3. Variables must have descriptions. These are automatically included in any published modules' auto-generated documentation through terraform-docs. Descriptions add additional context for new developers that descriptive names cannot.
4. Variables should have defined types.
5. Variables with non-environment-specific values (like disk size) should be given default values.
6. Variables for environment-specific values (like project\_id) should not be given defaults. This forces the calling module to provide meaningful values.

7. Variables should only have empty defaults (like empty strings or lists) where leaving the variable empty is a valid preference which will not be rejected by the underlying API(s).
8. Be thoughtful in your use of static literals (hardcoded strings, etc.) and parameterize anything which must vary per instance or environment. Local values can be used in cases where a literal is reused in multiple places without exposing it as a variable.
9. When deciding whether to expose a variable, ensure that you have a concrete use case for changing that variable. Don't expose variables on the off chance that it's needed.
  - Adding a variable with a default value is backwards compatible, and thus "cheap."
  - Removing a variable is backwards incompatible, and thus "expensive."
10. Boolean variables should be explicitly declared as strings until Terraform adds further support for boolean input variables.

#### **1.3.4 Outputs**

1. All outputs shall be organized into `outputs.tf`.
2. Outputs should have meaningful descriptions.
3. Output descriptions should be documented in the README. Descriptions should also be auto-generated on commit with `terraform-docs`.
4. Make an effort to output all the useful values root modules would want to reference or share with modules. Particularly for open source or heavily used modules, expose all outputs that have potential for consumption.

#### **1.3.5 Data Sources**

1. Data sources are located adjacent to the resources which reference them.
  - If you are fetching an image to be used in launching an instance, you can place it alongside the instance instead of collecting data resources in their own file.
2. If the number of data sources grows considerably, it's a reasonable practice to move these to a dedicated `data.tf` file.
3. Data sources should use variable or resource interpolation, where appropriate, to fetch data relative to your current environment.

#### **1.3.6 Scripts (Called by Terraform)**

1. Bespoke scripts can be called by Terraform through provisioners, including the `local-exec` provisioner.

2. Custom scripts should be avoided, if possible, and constrained to instances where native Terraform resources do not support the desired behavior. Any custom scripts used must have a clearly documented reasoning and ideally a deprecation plan. Additionally, custom scripts should not replace configuration management tools in cases where they are more appropriate.
3. Bespoke scripts called by Terraform must be organized into `scripts/`.
4. Use scripts only when absolutely necessary, as the state of resources created through scripts is not accounted for or managed by Terraform. You'll likely want to add a policy of `ignore_changes = [*]` on such resources.

### **1.3.7 Helper Scripts (not called by Terraform)**

1. Helper scripts should be organized in a `./helpers` directory.
2. Helper scripts should be documented in the README with an explanation and example invocations.
3. Helper scripts accepting arguments should provide argument-checking and `--help` output.

### **1.3.8 Static Files**

1. Static files which are referenced by Terraform (like Startup scripts loaded onto instances) but not executed must be organized into `files/`.
2. Lengthy heredocs should be externalized from their HCL and into external files. These should be referenced with the `file()` function.

### **1.3.9 Templates**

1. Files which are injected with the Terraform `template_file` resource should be given the file extension `.tpl`.
2. Templates must be placed in `templates/`.

### **1.3.10 Resources**

1. Resources that are the only one of their type (i.e., a single load balancer for an entire module) should be named 'main' to simplify references to that resource. It takes extra mental work to remember: `some_google_resource.my_unique_name.id` vs. `some_google_resource.main.id`
2. Resources that share the same type as others in the same module should be given meaningful names to differentiate them.

3. Resources must be named in snake-case (like db\_instance).
4. Resource names should be singular.
5. Resource names shouldn't repeat the resource type within the name. For example: Do this: `resource "google_compute_global_address" "main" { ... }` Not this: `resource "google_compute_global_address" "main_global_address" { ... }`
6. Ensure that deletion protection is enabled for stateful resources like databases. For example:

```
resource "google_sql_database_instance" "main" { name = "master-instance"
settings {
tier = "D0" }
lifecycle {
prevent_destroy = true
}
}
```

### **1.3.11 Formatting**

1. All Terraform files must conform to the standards of `terraform fmt`.

### **1.3.12 Expressions**

1. Limit the complexity of any individual interpolated expressions. If many functions are needed in a single expression, consider splitting it out into multiple expressions using locals.
2. Never have more than one ternary operation in a single line. Instead, use multiple local values to build up the logic.
3. Be sparing when using user-specified variables to set the `count` variable for resources. If a resource attribute is provided for such a variable (like `project_id`) and that resource does not yet exist, Terraform will not be able to generate a plan and will report the error “value of count cannot be computed.”
4. Use `count` to instantiate a resource conditionally. For example:

```
variable "readers" {
description = "..."
type = "list"
default = []
}
```

```
resource "foo" "bar" {
    // Do not create this resource if the list of readers is empty.
    count = "${length(var.readers) == 0 ? 0 : 1}"
    ...
}
```

## 1.4 Common Modules

Modules that are meant for reuse should follow the following standards, as well as the normal Terraform guidelines.

### 1.4.1 Structure

1. All common modules should have an OWNERSfile (or CODEOWNERS on GitHub) documenting who is responsible for the module.
2. Common modules should follow SemVer v2.0.0 when new versions are tagged/released.
3. Modules must not declare providers or backends. Leave that to the root modules.
  - Working examples should codify if a specific provider version is needed for a given module.

### 1.4.2 Variables

It's a good practice to allow flexibility in the labelling of resources through the module's interface. Consider providing a `labels` variable with a default value of an empty map to apply throughout labelable resources:

```
variable "labels" {
description = "A map of labels to apply to contained
resources."
default = {} type = "map"
}
```

### 1.4.3 Outputs

Outputs are required for common modules that define resources.

- Variables and outputs are used to infer dependencies between modules and resources. Without any outputs, users cannot properly order your module in relation to their Terraform configurations.

- Every resource defined in a common module should have at least one output which references that resource.

#### 1.4.4 Inline modules

1. Inline modules may be used to organize complex Terraform modules into smaller units, or de-duplicate common resources.
2. Inline modules shall be placed in `modules/$modulename`.
3. Inline modules should be treated as private and should not be used by outside modules, unless the common module specifically documents them otherwise.
4. Be aware that Terraform doesn't track refactored resources; if you start out with a number of resources in the top level module and then push them into submodules, Terraform will try to recreate all refactored resources.
5. Outputs defined by internal modules are not automatically exposed; if you want to share outputs from internal modules you'll need to re-output them.

### 1.5 Root Configs

Root configs, or root modules, are the working directories from which you run the Terraform CLI. They should follow the following standards, as well as the normal Terraform guidelines where applicable. Explicit recommendations for root modules supersede the general guidelines. Resources for different applications and projects should be separated into their own Terraform directories that can be managed independently of each other. A service might represent a particular application or a common service like shared networking. Importantly, all the Terraform code for a particular service should be nested under one directory (including subdirectories).

#### 1.5.1 Directory Structure

There are multiple ways to organize Terraform root configurations, especially when it comes to managing multiple environments. When it comes to managing the Terraform config for a particular service, the recommended structure is to use environment directories.

**Directories per environment** In this style, each service must split its Terraform config into multiple directories. In this structure, the directory layout must be as follows:

```
-- SERVICE-DIRECTORY/  
  -- OWNERS
```

```
-- modules/
  -- service/
    -- main.tf
    -- variables.tf
    -- outputs.tf
    -- README
-- ...other... -- environments/
  -- dev/
    -- backend.tf
    -- main.tf
    -- provider.tf
  -- qa/
    -- backend.tf
    -- main.tf
    -- provider.tf
  -- prod/
    -- backend.tf
    -- main.tf
    -- provider.tf
```

**Environment directories** Each environment directory within corresponds to a Terraform Workspace and deploys a version of the service to that environment. This config should reference modules to share code across environments, including typically a service module which includes the base shared Terraform config for the service.

This environment directory must contain the following files:

- \* A `backend.tf` file declaring the Terraform backend state location (typically GCS).
- \* A `main.tf` file which instantiates the service module.
- \* A `provider.tf` file which declares provider configuration.

***Workspaces per environment*** Alternatively, a single Terraform directory can be used per service and shared across environments. Each environment would have its own workspace. When using workspaces, all environments share the same modules, and the configuration is driven by a `tfvars` file and a workspace. Workspaces are helpful in that they limit the amount of code that must be copy-pasted between environment directories, which can help enforce parity between environments while maintaining their own state files. By default a single workspace named “`default`” exists. It is recommended to create and use a workspace for each environment and use the “`default`” workspace only when working with resources that may be used across multiple environments like some service accounts.

### 1.5.2 Outputs

1. Information from a root module which other root modules may depend on must be exported as outputs.
2. Root module outputs can be referenced using remote state.

#### ***Publishing outputs with remote states***

1. Make sure to re-output nested module outputs that are useful as remote state. Only root module-level outputs can be referenced from other Terraform environments/applications.
2. Information related to a service's endpoints should be exported to remote state to allow use by other dependent apps for configuration.

## 1.6 Versioning

### 1.6.1 Terraform

Terraform v0.12 is a significant release that will include some backwards incompatibilities. Pin the Terraform version (example given below) to a known safe version until v0.12 has been released and stabilized.

```
terraform {  
  required_version = "~> 0.11.10"  
}
```

### 1.6.2 Providers

Pin any providers to a known good version, and make updating the version pin a regular practice.

```
provider "google" { version = "~> 1.19.1"  
}
```

### 1.6.3 Modules

References to shared modules must be constrained to a release tag. Targeting a specific commit hash or branch is dangerous as it gives no context to the version of the underlying module. Updating modules should involve as little guesswork as possible for both authors and reviewers.

### 1.6.4 Constrain by git reference

References to shared modules may be constrained to any arbitrary git reference (commit, branch, or tag). For reasons outlined above, we only recommend using this to reference tags:

```
module "vpc" {  
    source = "git::https://github.com/terraform-google-modules/terraform-  
    google-network?ref=v0.4.0"  
    ... }
```

### 1.6.5 Constrain by version

When a git tag is released to the Terraform Module Registry, it creates a numbered version of that module (note that this does not apply to Github repositories, only to modules released to registries). An invocation can be constrained to said version:

```
module "nat_gateway" {  
    source = "GoogleCloudPlatform/nat-gateway/google"  
    version = "1.2.2"  
    ...  
}
```

# Exploring Modules and Testing

# Hello!

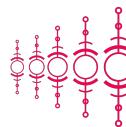
## Exploring Modules and Testing

Don't Repeat Yourself (DRY) is a key aspect of programming for applications and infrastructure. Just as important is testing your code works as intended.

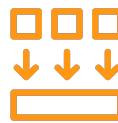




# DRY (Don't Repeat Yourself)



Reduce redundancy



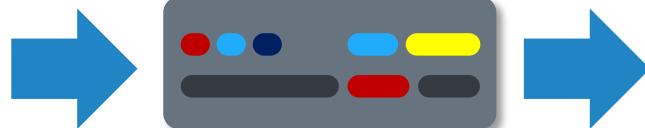
Create abstractions



Increase reusability

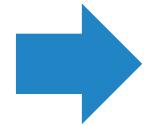
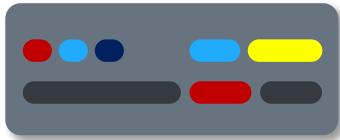
Fruit
Grapes
Apples
Oranges
Bananas
Pears

sorting\_script.sh



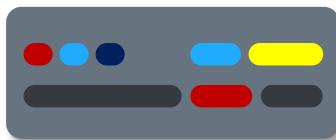
Fruit
Apples
Bananas
Grapes
Oranges
Pears

sorting\_script.sh

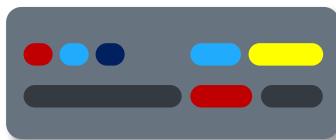


```
function sort(items)
{
    # Code to sort items
    ...
    ...
    return sorted_items;
}
```

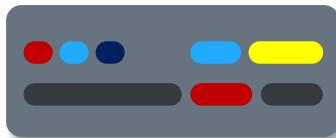
bubble\_sort



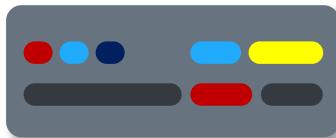
quick\_sort



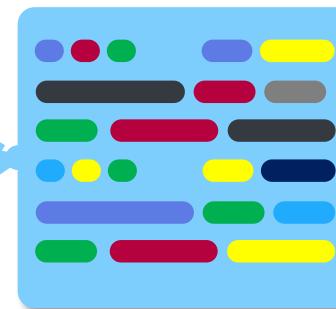
top\_sort



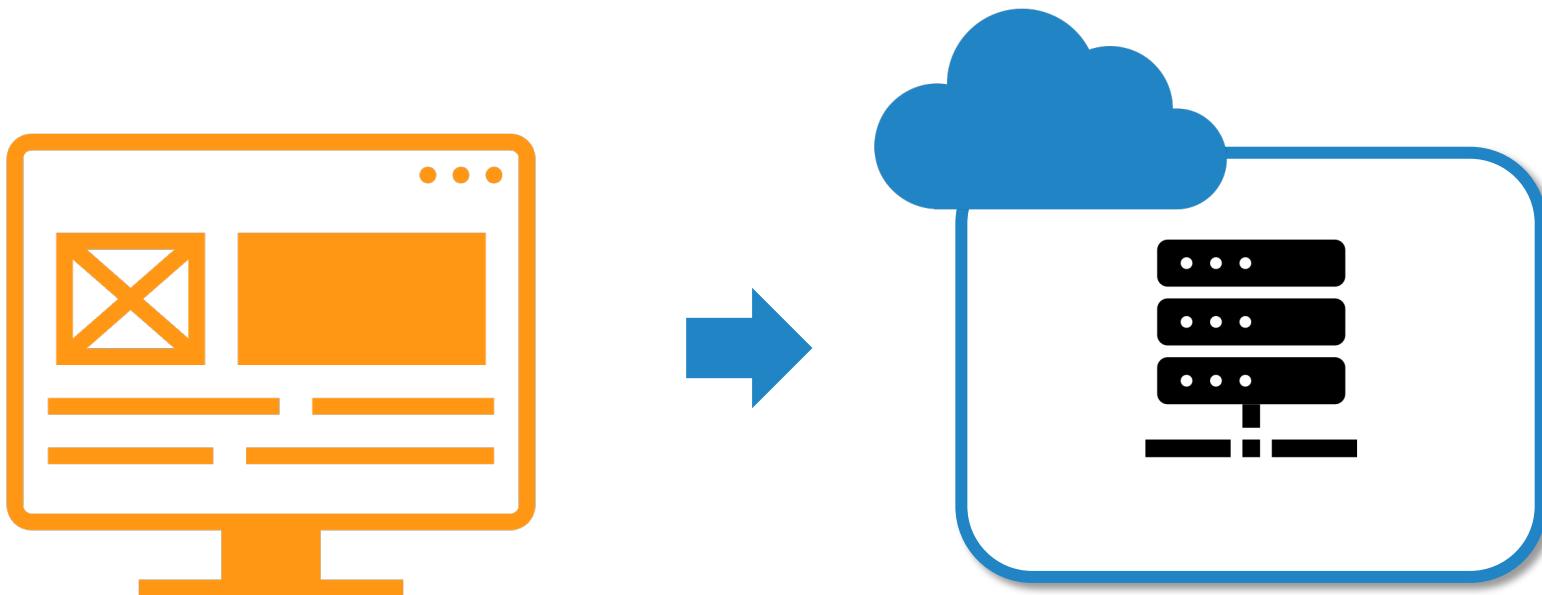
replacement\_sort

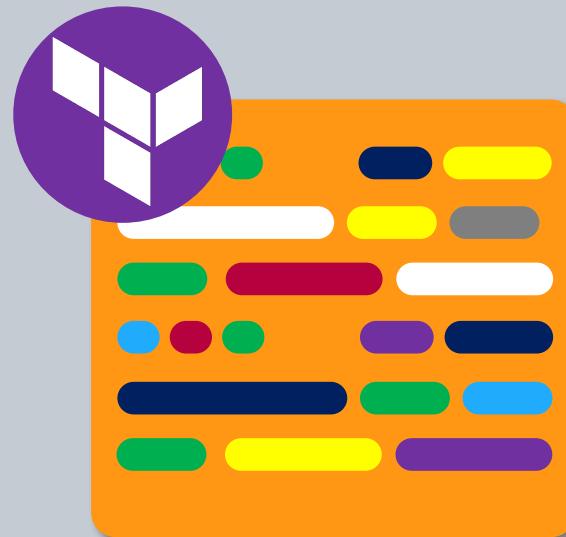


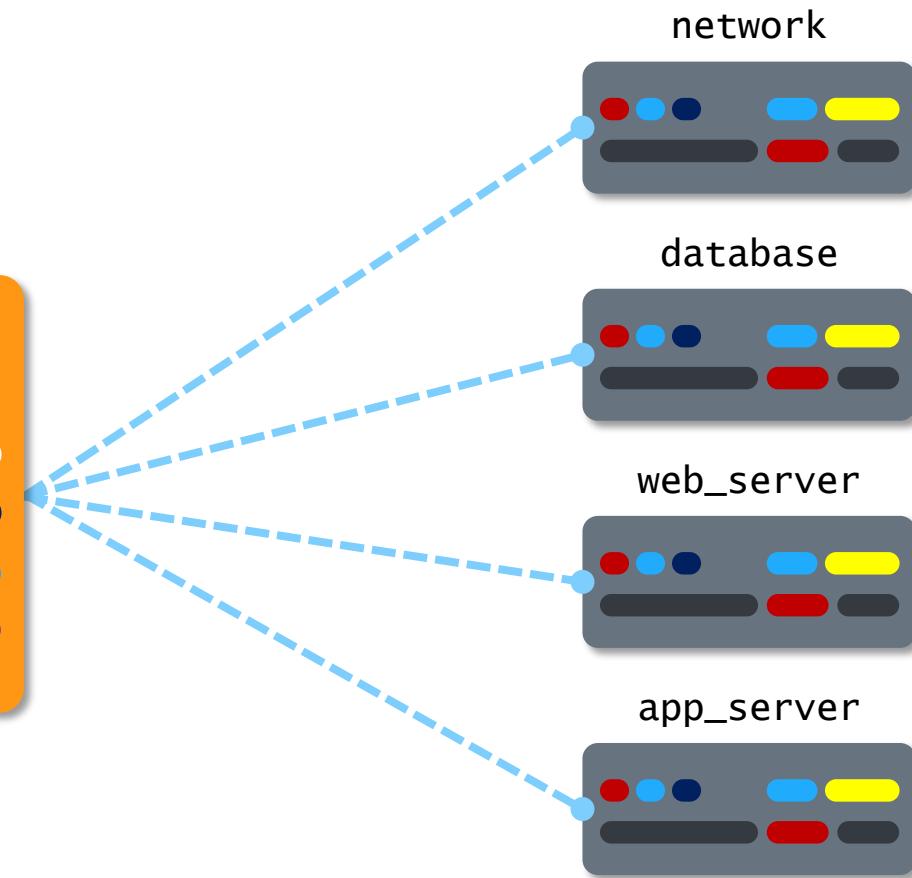
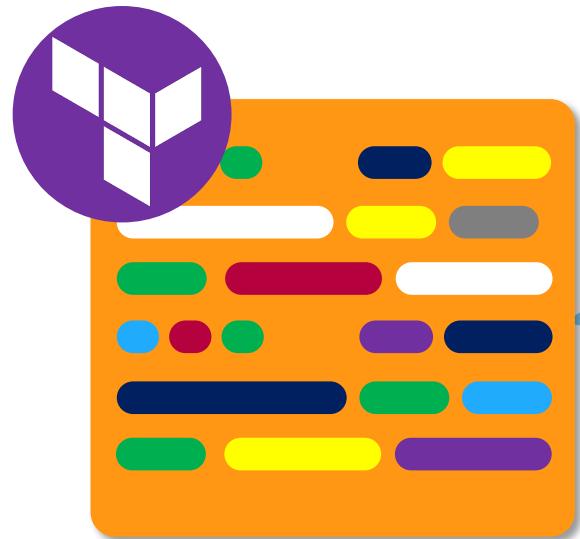
sorting.lib

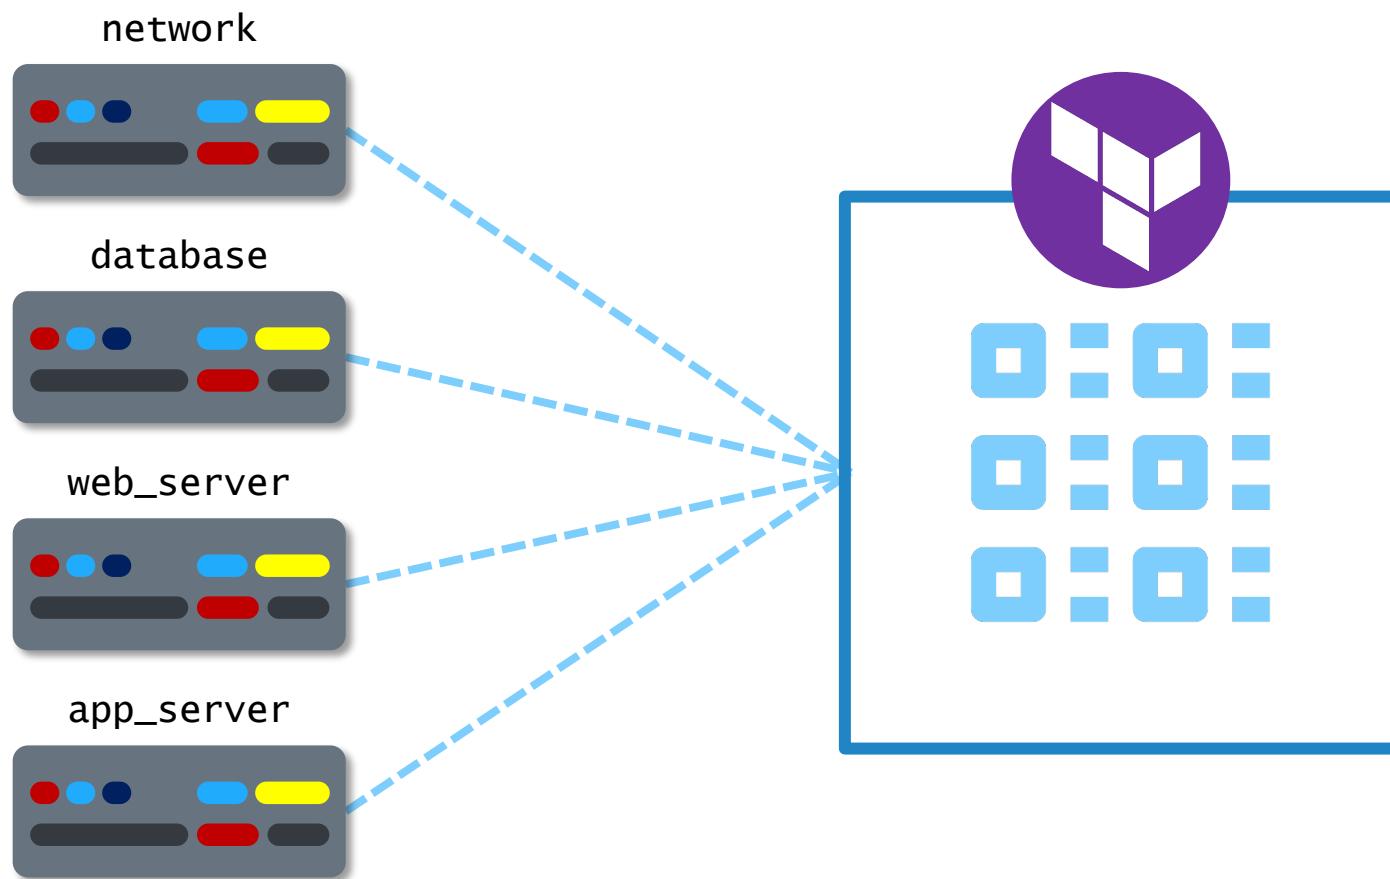


# Infrastructure Deployment











# Testing



Maintenance of code

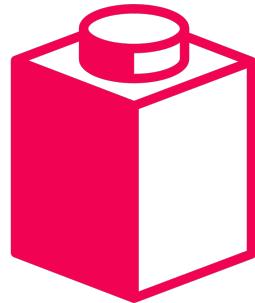


Breaking changes

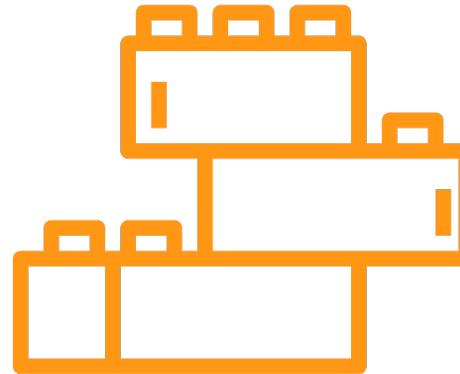


Use by others

# Test Types

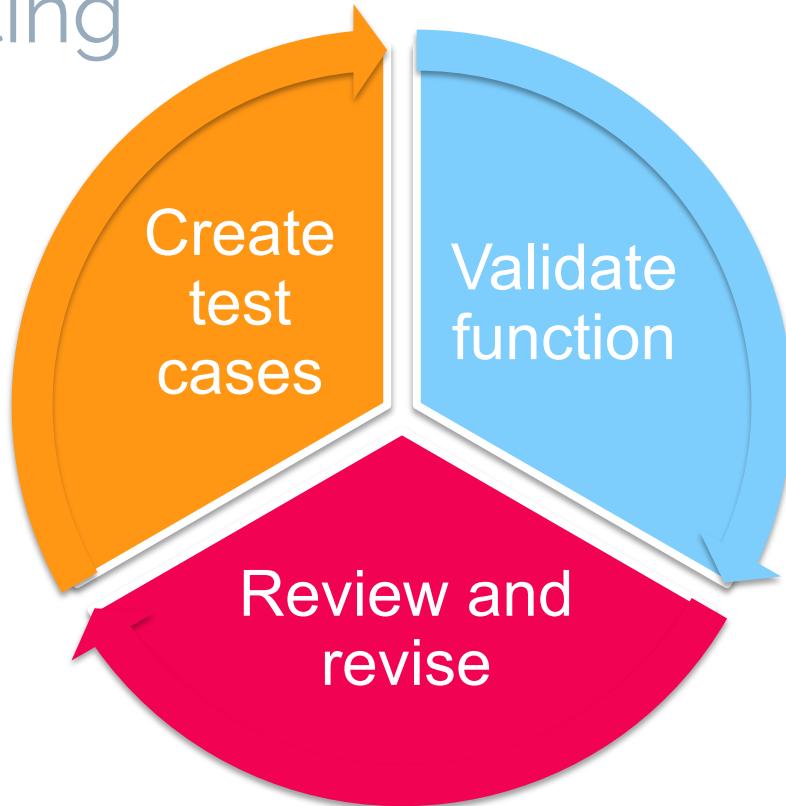


Unit testing



Integration testing

# Unit Testing



# Integration Testing



Units work together



Longer than unit test



Test in parallel



Retry on flakes

# Infrastructure as Code



Actual resources deployed for testing



Time consuming for large configurations



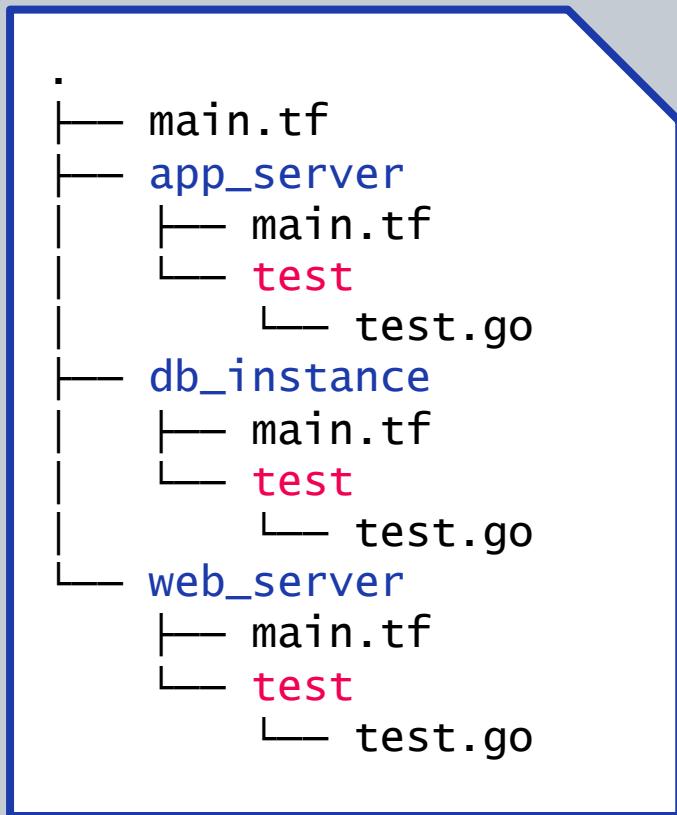
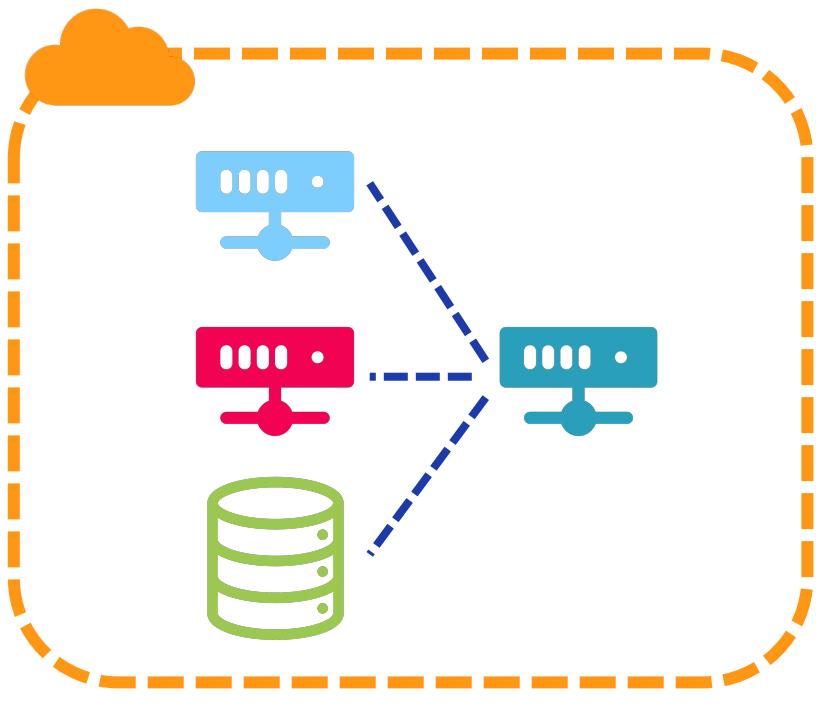
Unit testing handled by providers



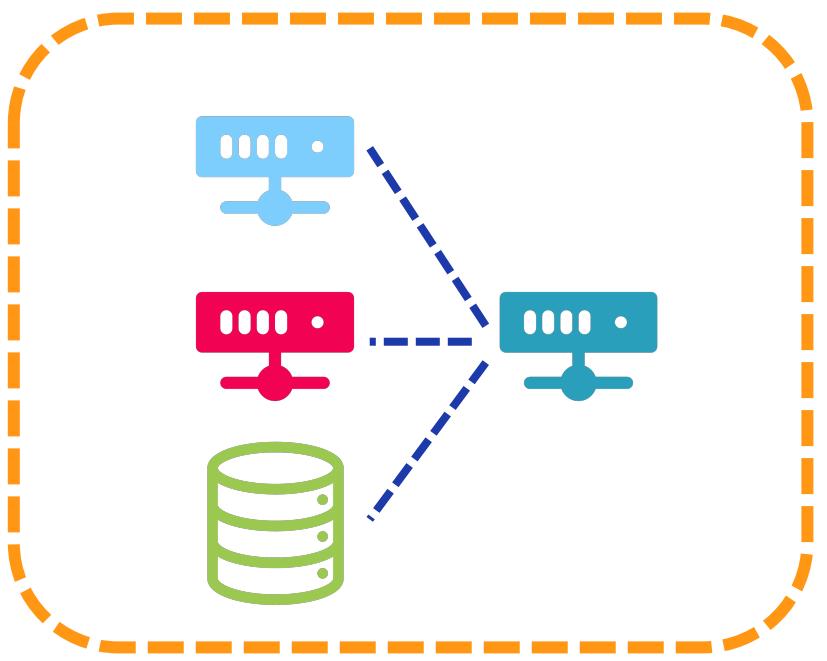
Test desired functionality of infrastructure



# Module Testing

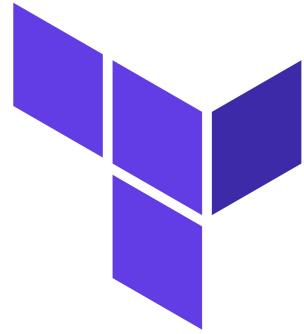


# Module Testing



```
web_server:  
- port: 443  
- status: 200  
  
application_server:  
- port: 443  
  
database_instance:  
- port: 3306  
- query: master_db
```

# Testing Tools



Terragrunt



Molecule



Checkov

# End-to-End Testing



Test everything together



Dedicated environment



Incremental changes



Validate changes only

# Summary

- ▶ Practice DRY when writing code for apps or infrastructure
- ▶ Infrastructure can be abstracted with modules and playbooks
- ▶ Testing is critical to dependable code
- ▶ IaC testing requires more resources and time





# API Driven Workflows

# Hello!

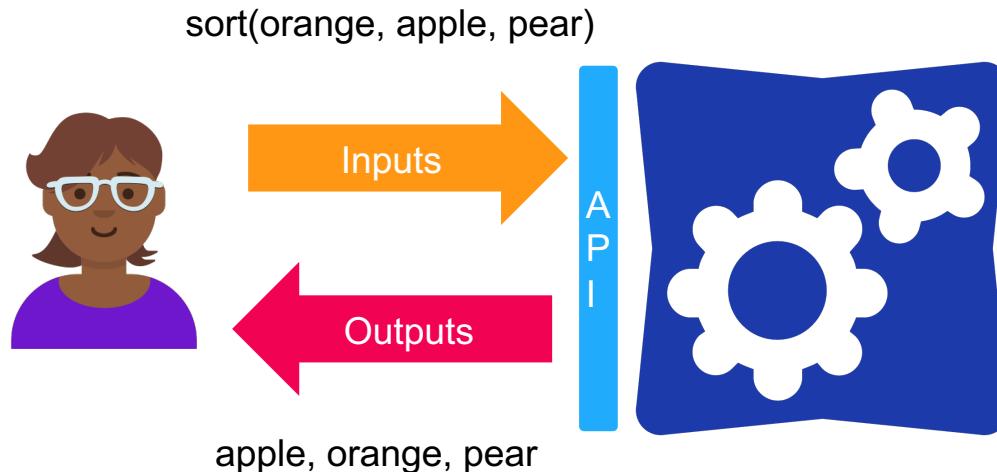
## API Driven Workflows

Orchestrating the process of creating and maintaining infrastructure by leveraging APIs to perform actions.

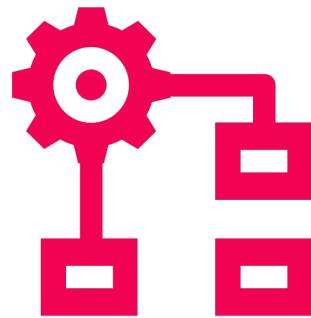


An Application Programming Interface is a way to interact with software using well-defined instructions. Using APIs enables the automation of manual tasks.

# Application Programming Interfaces (APIs)



# Types of APIs



RPC



REST



GraphQL

# API Interaction



Direct access



Graphical Interface



Precompiled tools



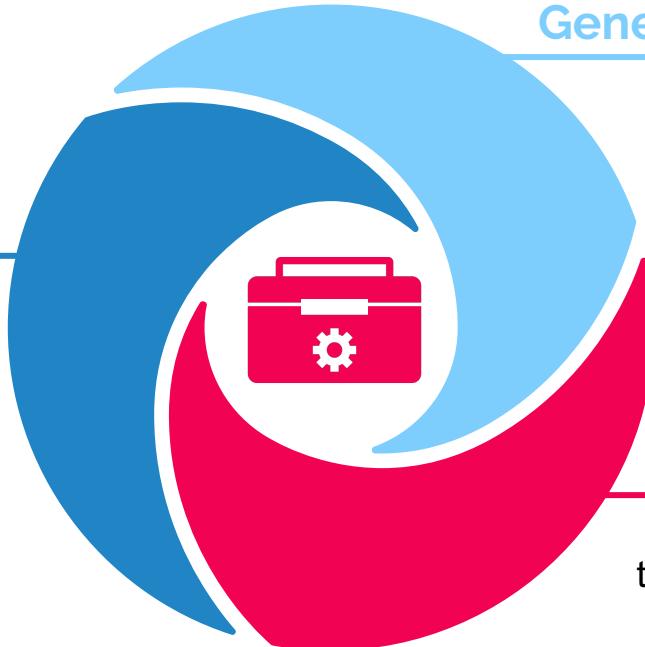
Software development kits



# Software Development Kits

## Use Common Tools

Leverage all the common programming tools like testing, linting, and scanning.



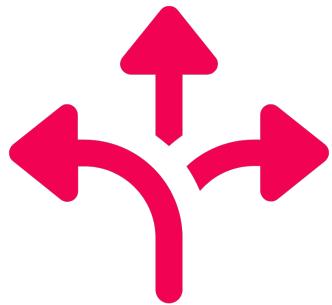
## General Purpose Language

SDKs provide libraries for general purpose programming languages to use.

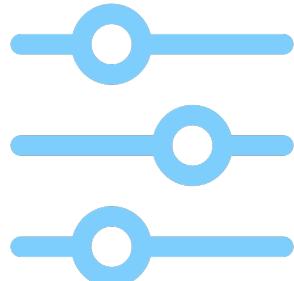
## Multiple APIs

A single SDK can stitch together multiple APIs into a unified library.

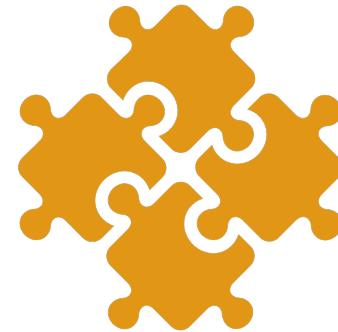
# API Driven Workflow Benefits



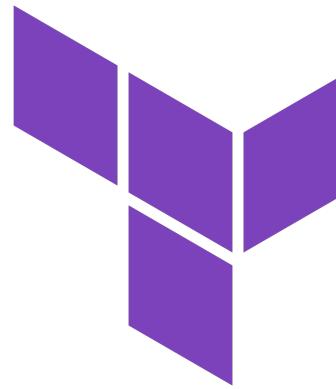
Flexibility



Customization

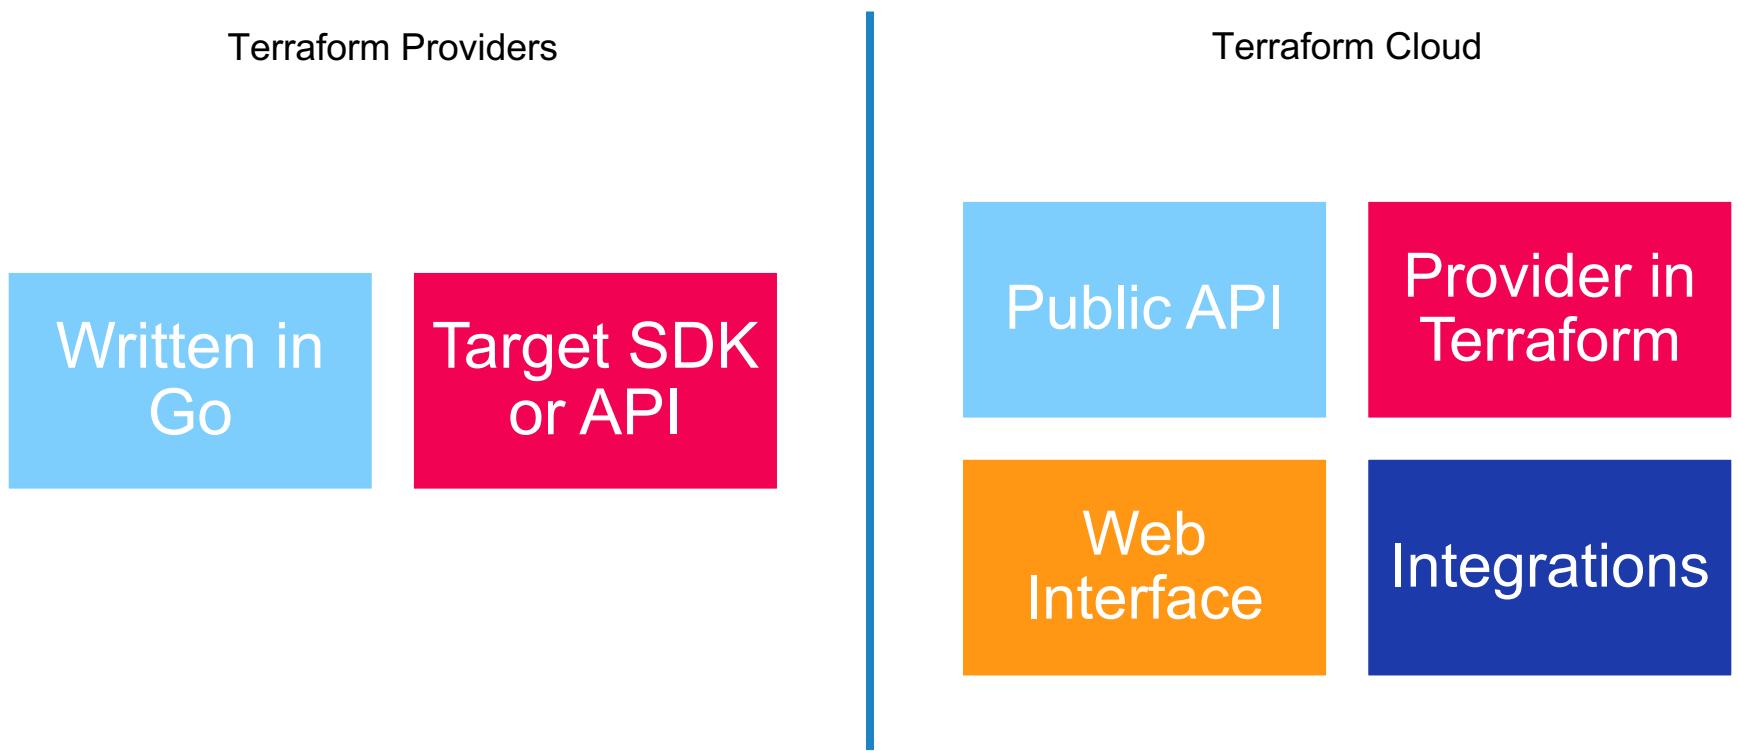


Integration



HashiCorp  
**Terraform**

# APIs in Terraform



# Terraform Cloud Workflows



Command line interface

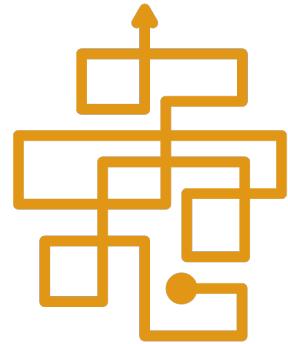


Version control system



Application programming interface

# API Workflow



Most complicated and  
flexible

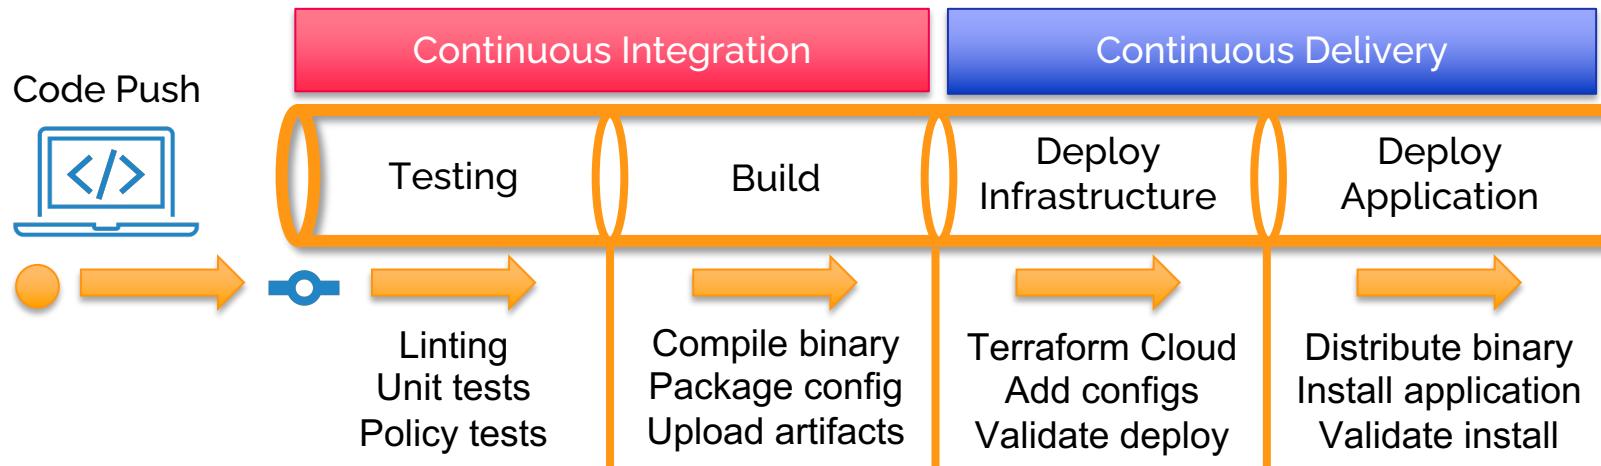


Bundle configurations  
in tarball



Relies on external  
orchestration

# API Workflow



# Summary

- ▶ Application programming interfaces are a contract between the client and the app
- ▶ APIs can be used in an automation workflow to interact with other services and applications
- ▶ Terraform Cloud and Enterprise have APIs that can integrate with your existing workflow

